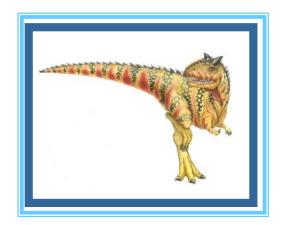
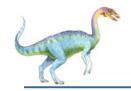
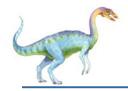
Chapter 1: Introduction





Syllabus

- Instructor
 - Yusung Wu (E-mail: hankwu@g2.nctu.edu.tw)
 - TAs: will announce later on E3
- Grading
 - Homework: 40%
 - Midterm: 30%
 - Final: 30%
- Textbook
 - Operating System Concepts, by Abraham Silberschatz, Greg Gagne, Peter Baer Galvin (8th or 9th Edition)
- What you will learn
 - virtual memory, threads, context switches, kernels, interrupts, system calls, interprocess communication, coordination, and the interaction between software and hardware

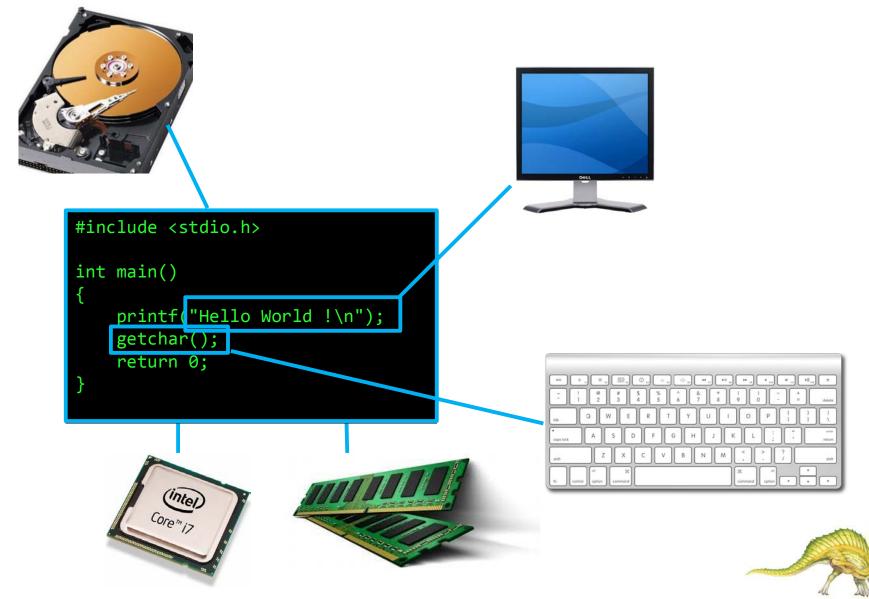


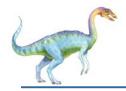
Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Structure
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management
- Protection and Security
- Distributed Systems
- Special-Purpose Systems
- Computing Environments
- Open-Source Operating Systems









- A program that acts as an intermediary between an application and the computer hardware
- A program that provides a convenient interface to the user

Application

```
#include <stdio.h>
int main()
{
    printf("Hello World !\n");
    getchar();
    return 0;
}
```

Operating System

Hardware

User





Controls and coordinates use of hardware among various applications

Application #1

```
#include <stdio.h>
int main()
{
    printf("Hello World 1\n");
    getchar();
    return 0;
}
```

Application #2

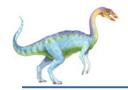
```
#include <stdio.h>

int main()
{
    printf("Hello World 2\n");
    getchar();
    return 0;
}
```

Operating System

Hardware

User



Provides a multi-user environment

Application #1

#include <stdio.h>

return 0;

```
int main()
{
    printf("Hello World !\n");
    getchar();
```

Application #2

```
#include <stdio.h>
int main()
{
    printf("Hello World !\n");
    getchar();
    return 0;
}
```

Operating System

Hardware

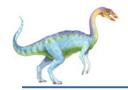
User #1

User #2

User #3

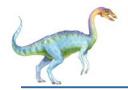
User #4





Distributed computing

```
Application #1
                                       Application #2
#include <stdio.h>
                                #include <stdio.h>
int main()
                                int main()
   printf("Hello World !\n");
                                    printf("Hello World !\n");
   getchar();
                                    getchar();
   return 0;
                                    return 0;
                    Operating System
      Hardware #1
                                      Hardware #2
 User #1
                 User #2
                                 User #3
                                                User #4
```



Cloud computing

Application #1

```
#include <stdio.h>
int main()
{
    printf("Hello World !\n");
    getchar();
    return 0;
}
```

Application #2

```
#include <stdio.h>
int main()
{
    printf("Hello World !\n");
    getchar();
    return 0;
}
```

Operating System

User #1

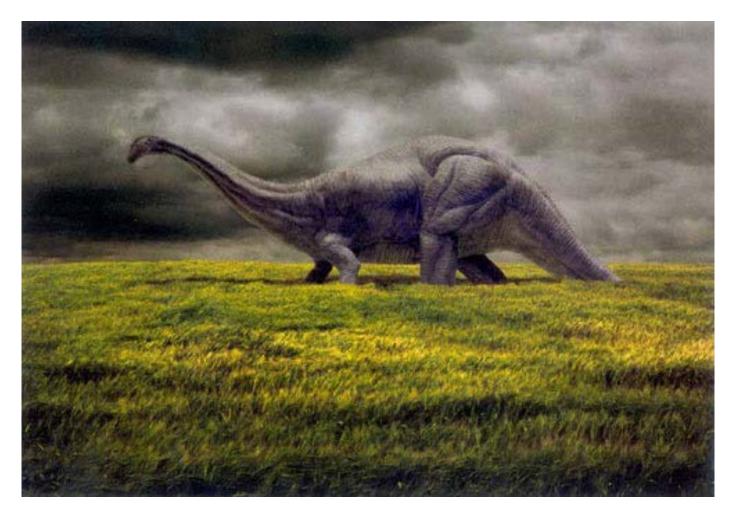
User #2

User #3

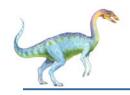
.

User #99,999,999





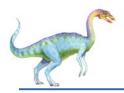




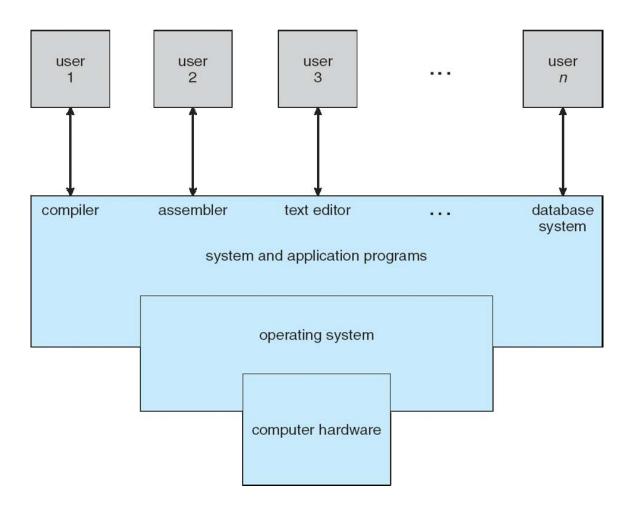
Computer System Structure

- Computer system can be divided into four components
 - Hardware provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - Operating system
 - Controls and coordinates use of hardware among various applications and users
 - Application programs define the ways in which the system resources are used to solve the computing problems of the users
 - Word processors, compilers, web browsers, database systems, video games
 - Users
 - People, machines, other computers





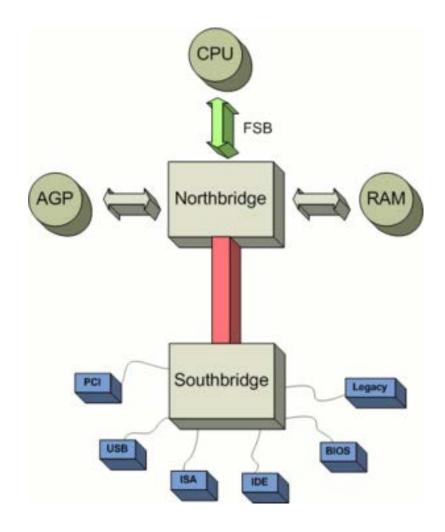
Four Components of a Computer System



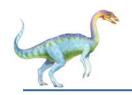




Computer System Hardware







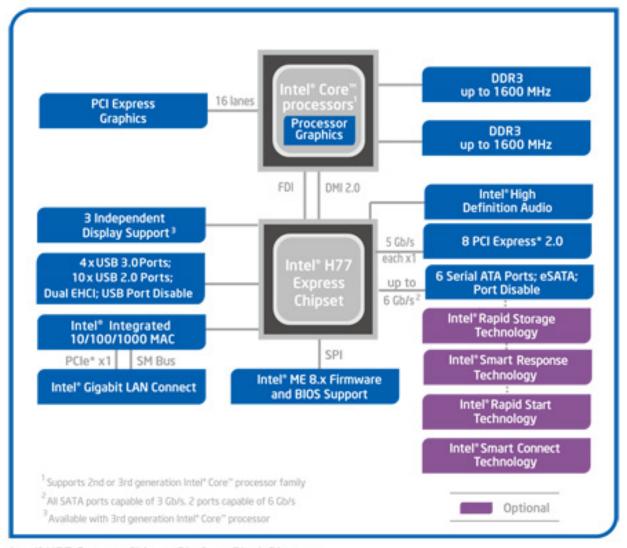
Computer System Hardware





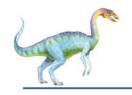


Computer Hardware









Operating System Definition

- OS is a piece of program (or pieces of programs)
- OS is a resource allocator
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a control program
 - Controls execution of programs to prevent errors and improper use of the computer





Operating System Definition (Cont)

- "Everything a vendor ships when you order an operating system" is good approximation
 - But varies wildly
- "The one program running at all times on the computer" is the kernel. Everything else is either a system program (ships with the operating system) or an application program

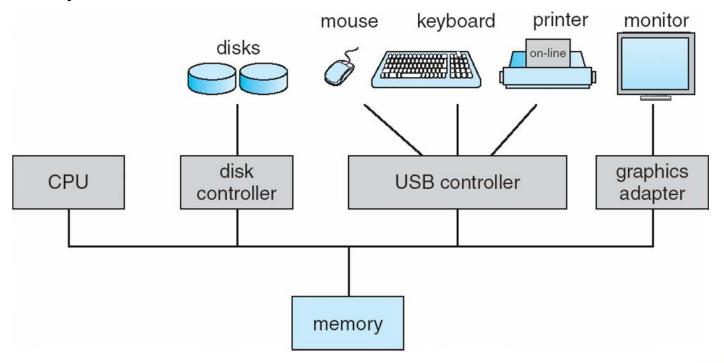
```
一命令提示字元
C:\Windows\System32>dir ntoskrnl.exe
 Volume in drive C has no label.
 Volume Serial Number is 90C4-5703
 Directory of C:\Windows\System32
2012/05/04 下午 07:06
                               5,559,664 ntoskrnl.exe
               1 File(s)
                              5,559,664 bytes
               0 Dir(s) 131,636,584,448 bytes free
C:\Windows\System32>copy ntoskrnl.exe d:\temp
       1 file(s) copied.
C:\Windows\System32>d:
D:\temp>dumpbin /disasm ntoskrnl.exe /out:ntoskrnl.asm
Microsoft (R) COFF/PE Dumper Version 10.00.40219.01
Copyright (C) Microsoft Corporation. All rights reserved.
D:\temp>
```

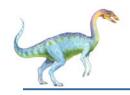




Computer System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles





Computer-System Operation

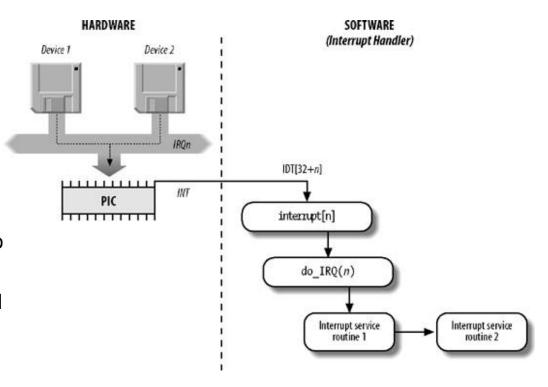
- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an interrupt





Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine generally, through the interrupt vector, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- Incoming interrupts are disabled while another interrupt is being processed to prevent a lost interrupt
- A trap is a software-generated interrupt caused either by an error or a user request
- An operating system is interrupt driven



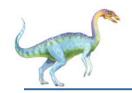




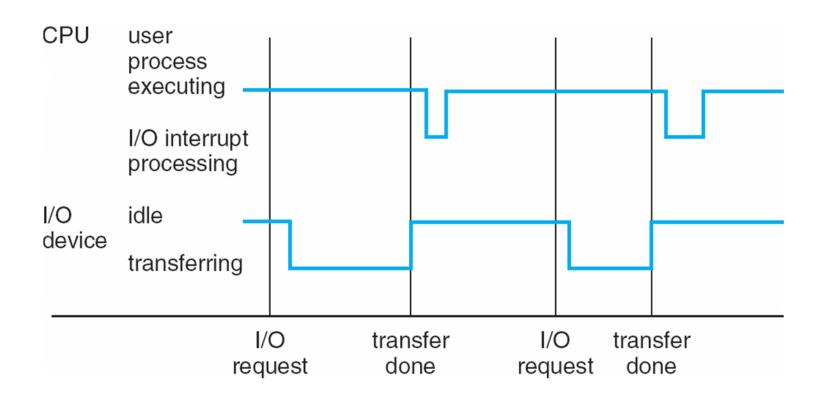
Interrupts

http://timetobleed.com/a-few-things-you-didnt-know-about-signals-in-linux-part-1/

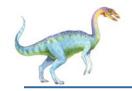
Interrupt vectors in Linux	
Vector range	Use
0-19 (0x0-0x13)	Nonmaskable interrupts and exceptions
20-31 (0x14-0x1f)	Intel-reserved
32-127 (0x20-0x7f)	External interrupts (IRQs)
128 (0x80)	Programmed exception for system calls
129-238 (0x81-0xee)	External interrupts (IRQs)
239 (0xef)	Local APIC timer interrupt
240 (0xf0)	Local APIC thermal interrupt (introduced in the Pentium 4 models)
241-250 (0xf1-0xfa)	Reserved by Linux for future use
251-253 (0xfb-0xfd)	Interprocessor interrupts
254 (0xfe)	Local APIC error interrupt (generated when the local APIC detects an erroneous condition)
255 (0xff)	Local APIC spurious interrupt (generated if the CPU masks an interrupt while the hardware device raises it)



Interrupt Timeline







I/O Structure

- After I/O starts, control returns to user program only upon I/O completion
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access)
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing
- [I/O multiplexing] After I/O starts, control returns to user program without waiting for I/O completion
 - System call request to the operating system to allow user to wait for I/O completion
 - Device-status table contains entry for each I/O device indicating its type, address, and state
 - Operating system indexes into I/O device table to determine device status and to modify table entry to include interrupt



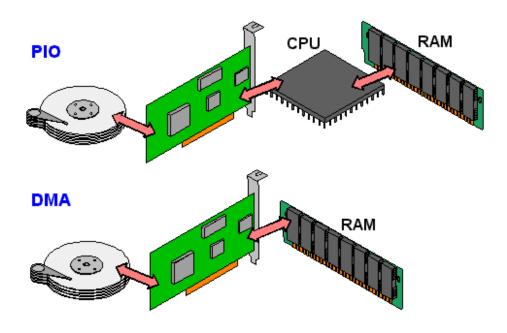


Direct Memory Access Structure

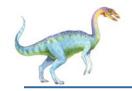
- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte

From Computer Desktop Encyclopedia

3 1998 The Computer Language Co. Inc.

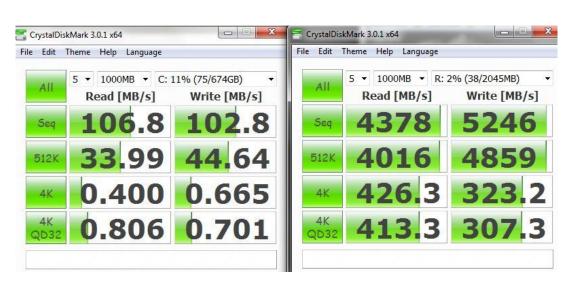


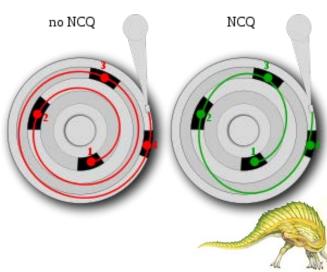


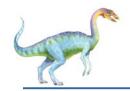


Storage Structure

- Main memory only large storage media that the CPU can access directly
- Secondary storage extension of main memory that provides large nonvolatile storage capacity
- Magnetic disks rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into tracks, which are subdivided into sectors
 - The disk controller determines the logical interaction between the device and the computer







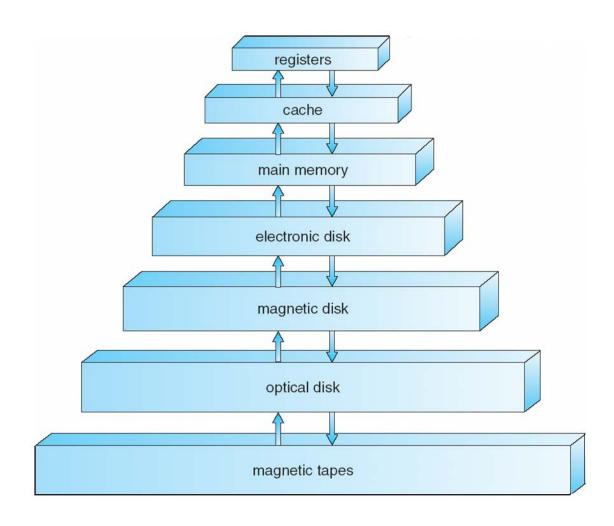
Storage Hierarchy

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- Caching copying information into faster storage system; main memory can be viewed as a last cache for secondary storage

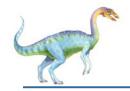




Storage-Device Hierarchy



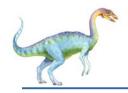




Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy





Computer-System Architecture

- Most systems use a single general-purpose processor (PDAs through mainframes)
 - Most systems have special-purpose processors as well
- Multiprocessors systems growing in use and importance
 - Also known as parallel systems, tightly-coupled systems
 - Advantages include
 - Increased throughput
 - 2. Economy of scale
 - 3. Increased reliability graceful degradation or fault tolerance
 - Two types
 - 1. Asymmetric Multiprocessing
 - 2. Symmetric Multiprocessing





Computer System Architecture

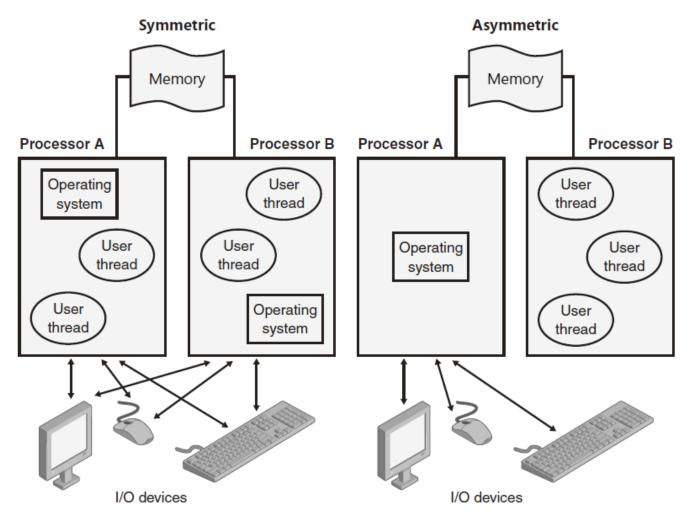
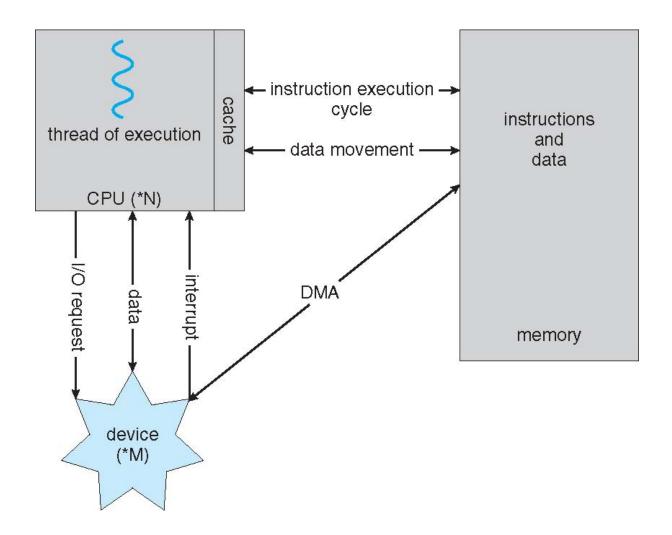


FIGURE 2-2 Symmetric vs. asymmetric multiprocessing





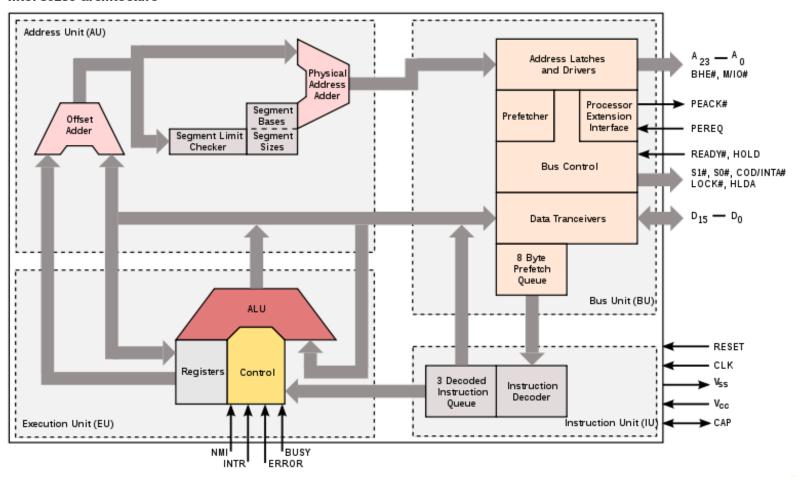
How a Modern Computer Works

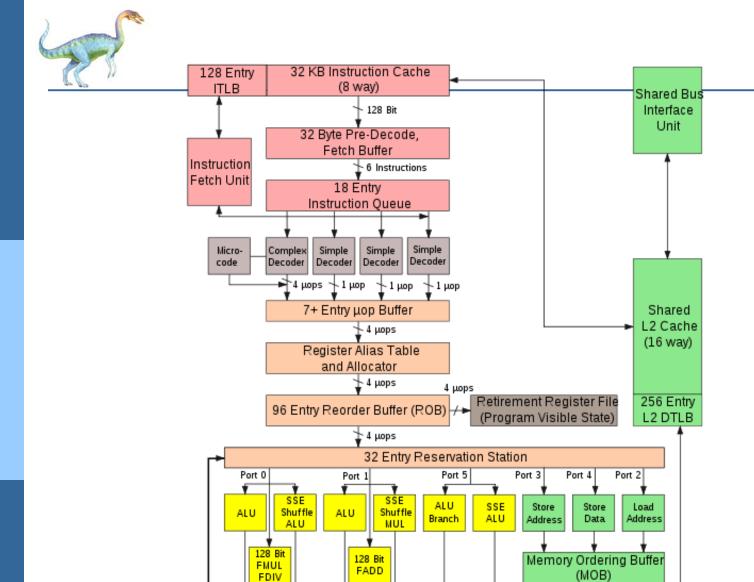






Intel 80286 architecture







Internal Results Bus

128 Bit

32 KB Dual Ported Data Cache

(8 way)

Store

128 Bit

Load

16 Entry

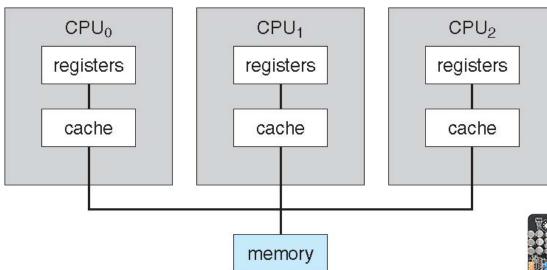
DTLB

256

Bit

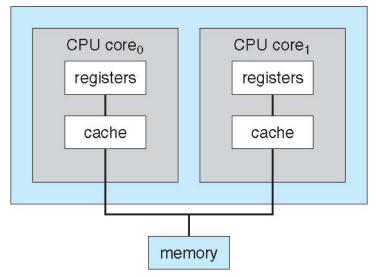


Symmetric Multiprocessing Architecture

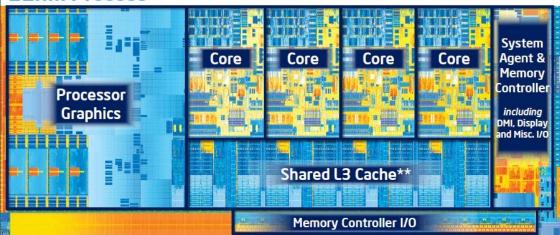


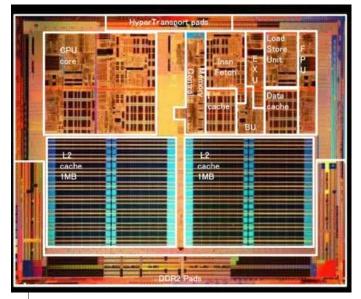


A Dual-Core Design



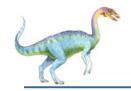






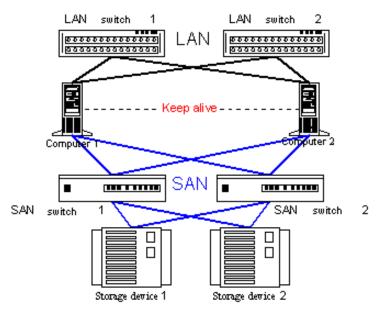
AMD 64 X2 5000+dual-core processor

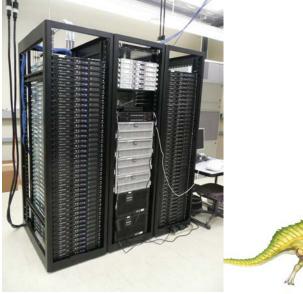




Clustered Systems

- Like multiprocessor systems, but multiple systems working together
 - Usually sharing storage via a storage-area network (SAN)
 - Provides high-availability services that survive hardware failures
 - Asymmetric clustering has one machine in hot-standby mode
 - Symmetric clustering has multiple nodes running applications, monitoring each other
 - Some clusters are for high-performance computing (HPC)
 - Applications must be written to use parallelization



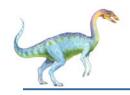




Computer Startup

- bootstrap program is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, generally known as firmware
 - Initializes all aspects of system
 - Loads operating system kernel and starts execution
- Operating system probes and initializes hardware
- Hardware provides the driving power
- Operating system distributes the power to system services and applications (jobs)



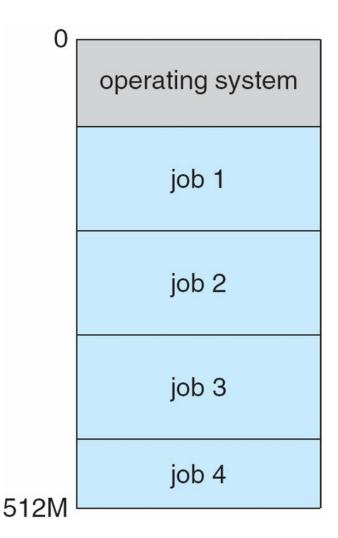


Operating System Structure

- Multiprogramming needed for efficiency
 - Single user cannot keep CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - A subset of total jobs in system is kept in memory
 - One job selected and run via job scheduling
 - When it has to wait (for I/O for example), OS switches to another job
- Timesharing (multitasking) is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating interactive computing
 - Response time should be < 1 second
 - Each user has at least one program executing in memory ⇒process
 - If several jobs ready to run at the same time ⇒ CPU scheduling
 - If processes don't fit in memory, swapping moves them in and out to run
 - Virtual memory allows execution of processes not completely in memory



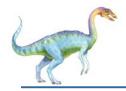
Memory Layout for Multiprogrammed System



Question:

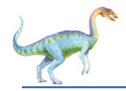
Should we allow job→OS control flow?





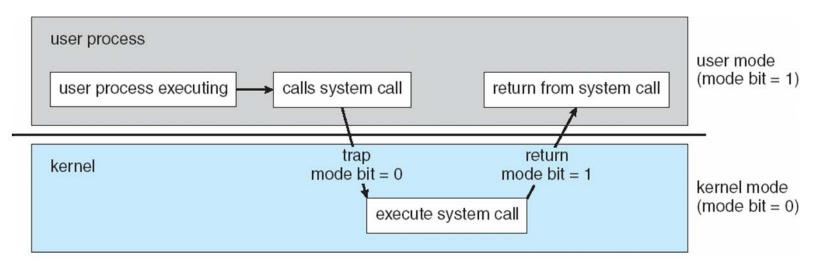
Operating-System Operations

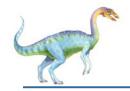
- Interrupt driven by hardware
- Software error creates exception
 - Division by zero,
- Software request creates trap
 - request for operating system service
- Other process problems include infinite loop, processes modifying each other or the operating system
- Dual-mode operation allows OS to protect itself and other system components
 - User mode and kernel mode
 - Mode bit provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as privileged, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user



Transition from User to Kernel Mode

- Timer to prevent infinite loop / process hogging resources
 - Set interrupt after specific period
 - Operating system decrements counter
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time





Process Management

- A process is a program in execution. It is a unit of work within the system. Program is a passive entity, process is an active entity.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one program counter specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes / threads



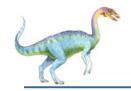


Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling

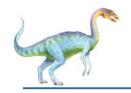




Memory Management

- All data in memory before and after processing
- All instructions in memory in order to execute
- Memory management determines what is in memory when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

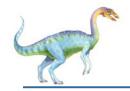




Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit file
 - Each medium is controlled by device (i.e., disk drive, tape drive)
 - Varying properties include access speed, capacity, datatransfer rate, access method (sequential or random)
- File-System management
 - Files usually organized into directories
 - Access control on most systems to determine who can access what
 - OS activities include
 - Creating and deleting files and directories
 - Primitives to manipulate files and dirs
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media

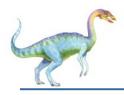




Mass-Storage Management

- Usually disks used to store data that does not fit in main memory or data that must be kept for a "long" period of time
- Proper management is of central importance
- Entire speed of computer operation hinges on disk subsystem and its algorithms
- OS activities
 - Free-space management
 - Storage allocation
 - Disk scheduling
- Some storage need not be fast
 - Tertiary storage includes optical storage, magnetic tape
 - Still must be managed
 - Varies between WORM (write-once, read-many-times) and RW (read-write)





Performance of Various Levels of Storage

Movement between levels of storage hierarchy can be explicit or implicit

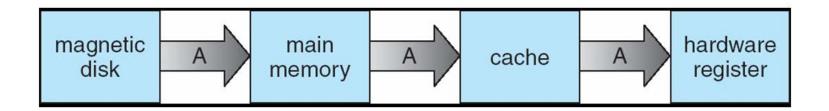
Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 - 100,000	5000 - 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape



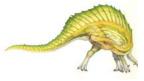


Migration of Integer A from Disk to Register

 Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide cache coherency in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex
 - Several copies of a datum can exist
 - Various solutions covered in Chapter 17

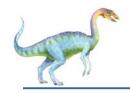




I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user
- I/O subsystem responsible for
 - Memory management of I/O including buffering (storing data temporarily while it is being transferred), caching (storing parts of data in faster storage for performance), spooling (the overlapping of output of one job with input of other jobs)
 - General device-driver interface
 - Drivers for specific hardware devices

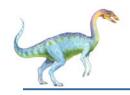




Protection and Security

- Protection any mechanism for controlling access of processes or users to resources defined by the OS
- Security defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine who can do what
 - User identities (user IDs, security IDs) include name and associated number, one per user
 - User ID then associated with all files, processes of that user to determine access control
 - Group identifier (group ID) allows set of users to be defined and controls managed, then also associated with each process, file
 - Privilege escalation allows user to change to effective ID with more rights

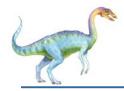




Computing Environments

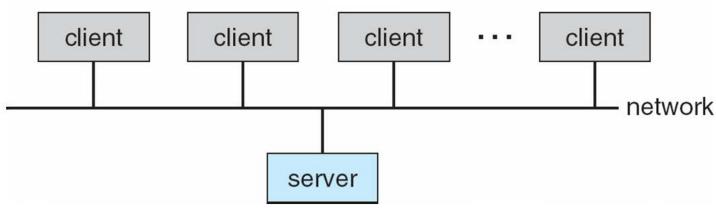
- Traditional computer
 - Blurring over time
 - Office environment
 - PCs connected to a network, terminals attached to mainframe or minicomputers providing batch and timesharing
 - Now portals allowing networked and remote systems access to same resources
 - Home networks
 - Used to be single system, then modems
 - Now firewalled, networked



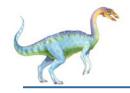


Computing Environments (Cont)

- Client-Server Computing
 - Dumb terminals supplanted by smart PCs
 - Many systems now servers, responding to requests generated by clients
 - Compute-server provides an interface to client to request services (i.e. database)
 - File-server provides interface for clients to store and retrieve files



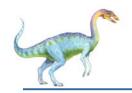




Peer-to-Peer Computing

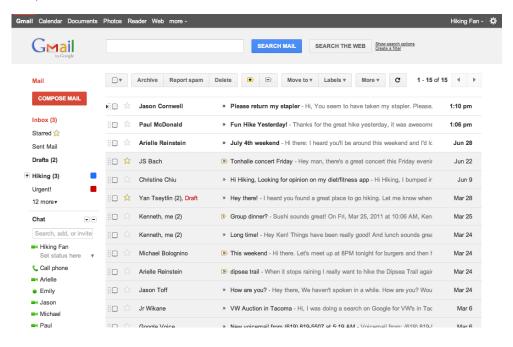
- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - Registers its service with central lookup service on network, or
 - Broadcast request for service and respond to requests for service via discovery protocol
 - Examples include Napster and Gnutella





Web-Based Computing

- Web has become ubiquitous
- PCs most prevalent devices
- More devices becoming networked to allow web access
- New category of devices to manage web traffic among similar servers: load balancers
- Use of operating systems like Windows 95, client-side, have evolved into Linux and Windows XP, which can be clients and servers







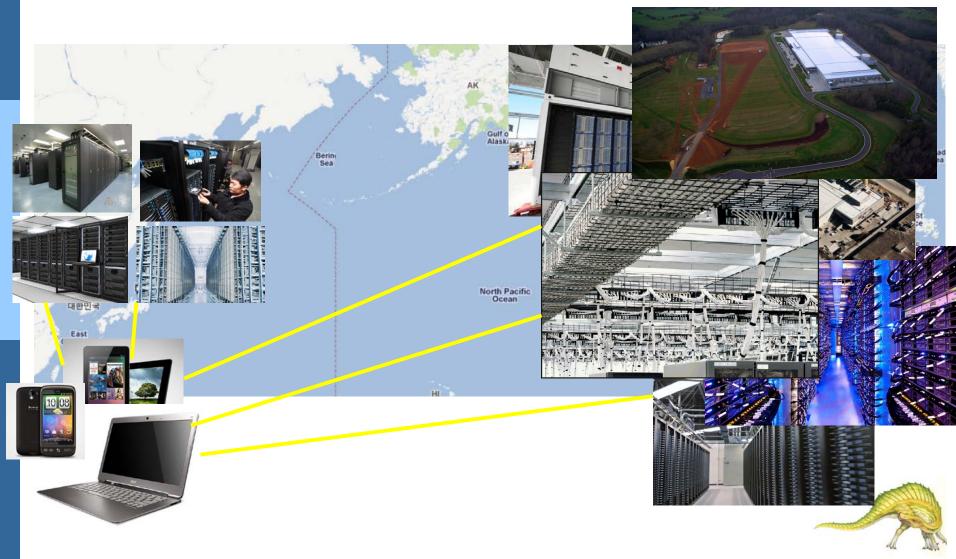
Mobile Computing

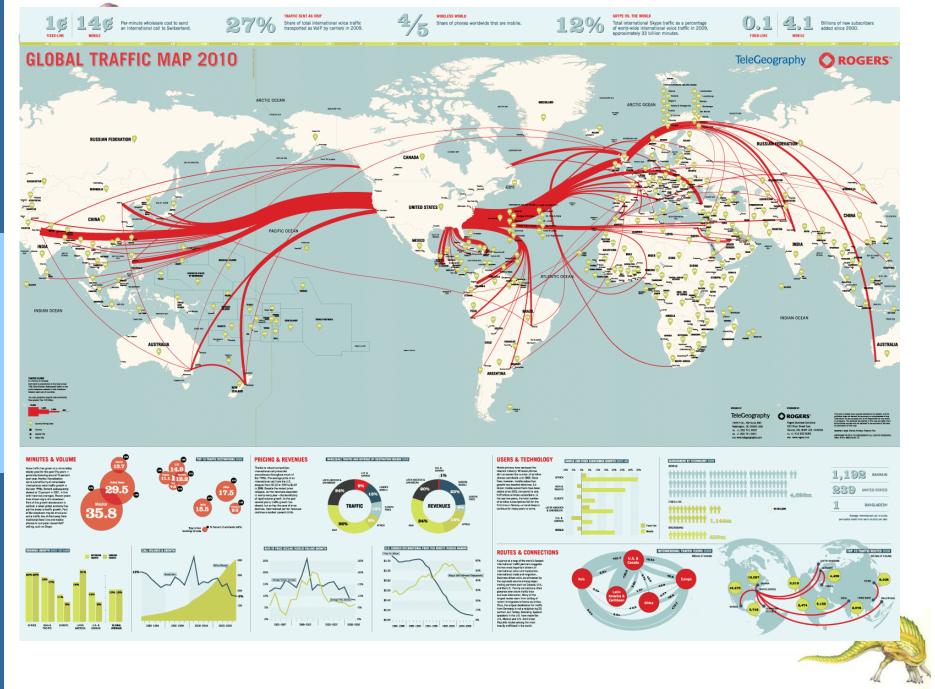






Cloud Computing







Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary closed-source
- Counter to the copy protection and Digital Rights Management (DRM)
 movement
- Started by Free Software Foundation (FSF), which has "copyleft" GNU Public License (GPL)
- Examples include GNU/Linux, BSD UNIX (including core of Mac OS X), and Sun Solaris



End of Chapter 1

