


# Chapter 5: Process Scheduling

---





# Chapter 5: Process Scheduling

---

- ☐ Basic Concepts
- ☐ Scheduling Criteria
- ☐ Scheduling Algorithms
- ☐ Thread Scheduling
- ☐ Multiple-Processor Scheduling
- ☐ Operating Systems Examples
- ☐ Algorithm Evaluation





# Process Scheduling

---

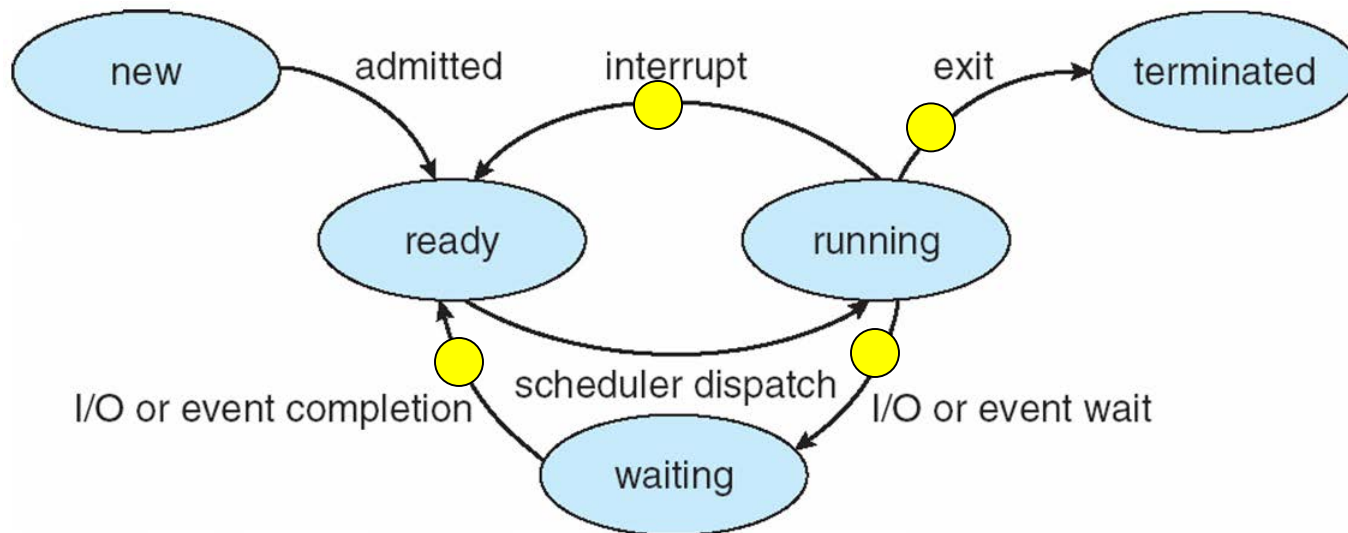
- ☐ Have M jobs ready to run
- ☐ Have  $N \geq 1$  CPUs
- ☐ Which job to assign to which CPU(s) at what time?





# CPU Scheduler

- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready
  4. Terminates
- **nonpreemptive scheduler uses 1 and 4**
- **Preemptive scheduler kicks in for all four time points**





# CPU Scheduler

```
LXR linux/fs/block_d x
lxr.linux.no/linux+v3.6.3/fs/block_dev.c#L748

711         else
712             return true;    /* is a partition of an un-held device */
713     }
714
715 /**
716  * bd_prepare_to_claim - prepare to claim a block device
717  * @bdev: block device of interest
718  * @whole: the whole device containing @bdev, may equal @bdev
719  * @holder: holder trying to claim @bdev
720  *
721  * Prepare to claim @bdev. This function fails if @bdev is already
722  * claimed by another holder and waits if another claiming is in
723  * progress. This function doesn't actually claim. On successful
724  * return, the caller has ownership of bd_claiming and bd_holder[s].
725  *
726  * CONTEXT:
727  * spin_lock(&bdev_lock). Might release bdev_lock, sleep and regrab
728  * it multiple times.
729  *
730  * RETURNS:
731  * 0 if @bdev can be claimed, -EBUSY otherwise.
732  */
733 static int bd_prepare_to_claim(struct block_device *bdev,
734                               struct block_device *whole, void *holder)
735 {
736     retry:
737         /* if someone else claimed, fail */
738         if (!bd_may_claim(bdev, whole, holder))
739             return -EBUSY;
740
741         /* if claiming is already in progress, wait for it to finish */
742         if (whole->bd_claiming) {
743             wait_queue_head_t *wq = bit_waitqueue(&whole->bd_claiming, 0);
744             DEFINE_WAIT(wait);
745
746             prepare_to_wait(wq, &wait, TASK_UNINTERRUPTIBLE);
747             spin_unlock(&bdev_lock);
748             schedule();
749             finish_wait(wq, &wait);
750             spin_lock(&bdev_lock);
751             goto retry;
752         }
753
754         /* yay, all mine */
755         return 0;
756     }
```





# Dispatcher

---

- ☐ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - ☐ switching context
  - ☐ switching to user mode
  - ☐ jumping to the proper location in the user program to restart that program
- ☐ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





# Context switch

```
← → C | lxr.linux.no/linux+v3.6.3/kernel/sched/core.c#L2046
2046 context_switch(struct rq *rq, struct task_struct *prev,
2047                 struct task_struct *next)
2048 {
2049     struct mm_struct *mm, *oldmm;
2050
2051     prepare_task_switch(rq, prev, next);
2052
2053     mm = next->mm;
2054     oldmm = prev->active_mm;
2055     /*
2056      * For paravirt, this is coupled with an exit in switch_to to
2057      * combine the page table reload and the switch backend into
2058      * one hypercall.
2059      */
2060     arch_start_context_switch(prev);
2061
2062     if (!mm) {
2063         next->active_mm = oldmm;
2064         atomic_inc(&oldmm->mm_count);
2065         enter_lazy_tlb(oldmm, next);
2066     } else
2067         switch_mm(oldmm, mm, next);
2068
2069     if (!prev->mm) {
2070         prev->active_mm = NULL;
2071         rq->prev_mm = oldmm;
2072     }
2073     /*
2074      * Since the runqueue lock will be released by the next
2075      * task (which is an invalid locking op but in the case
2076      * of the scheduler it's an obvious special-case), so we
2077      * do an early lockdep release here:
2078      */
2079 #ifndef __ARCH_WANT_UNLOCKED_CTXSW
2080     spin_release(&rq->lock.dep_map, 1, _THIS_IP_);
2081 #endif
2082
2083     /* Here we just switch the register state and the stack. */
2084     switch_to(prev, next, prev);
2085
2086     barrier();
2087     /*
2088      * this_rq must be evaluated again because prev may have moved
2089      * CPUs since it called schedule(), thus the 'rq' on its stack
2090      * frame will be invalid.
2091      */
2092     finish_task_switch(this_rq(), prev);
2093 }
```

switching  
address  
space

switching register  
state and stack





# Context switch

```
← → C | lxr.linux.no/linux+v3.6.3/arch/x86/include/asm/switch_to.h
80
81 /* frame pointer must be last for get_wchan */
82 #define SAVE_CONTEXT    "pushf ; pushq %%rbp ; movq %%rsi,%%rbp\n\t"
83 #define RESTORE_CONTEXT "movq %%rbp,%%rsi ; popq %%rbp ; popf\t"
84
85 #define EXTRA_CLOBBER \
86     , "rcx", "rbx", "rdx", "r8", "r9", "r10", "r11", \
87     "r12", "r13", "r14", "r15"
88
89 #ifdef CONFIG_CC_STACKPROTECTOR
90 #define switch_canary \
91     "movq %P[task_canary](%%rsi),%%r8\n\t" \
92     "movq %%r8, \"__percpu_arg([gs_canary])\" \n\t"
93 #define switch_canary_oparam \
94     , [gs_canary] "=m" (irq_stack_union.stack_canary)
95 #define switch_canary_iparam \
96     , [task_canary] "i" (offsetof(struct task_struct, stack_canary))
97 #else /* CC_STACKPROTECTOR */
98 #define switch_canary
99 #define switch_canary_oparam
100 #define switch_canary_iparam
101 #endif /* CC_STACKPROTECTOR */
102
103 /* Save restore flags to clear handle leaking NT */
104 #define switch_to(prev, next, last) \
105     asm volatile(SAVE_CONTEXT \
106         "movq %%rsp,%P[threadrsp](%[prev])\n\t" /* save RSP */ \
107         "movq %P[threadrsp](%[next]),%%rsp\n\t" /* restore RSP */ \
108         "call __switch_to\n\t" \
109         "movq \"__percpu_arg([current_task])\",%%rsi\n\t" \
110         switch_canary \
111         "movq %P[thread_info](%%rsi),%%r8\n\t" \
112         "movq %%rax,%%rdi\n\t" \
113         "testl %[tif_fork],%P[ti_flags](%%r8)\n\t" \
114         "jnz ret_from_fork\n\t" \
115         RESTORE_CONTEXT \
116         : "=a" (last) \
117         : \
118         : [next] "S" (next), [prev] "D" (prev), \
119         [threadrsp] "i" (offsetof(struct task_struct, thread.sp)), \
120         [ti_flags] "i" (offsetof(struct thread_info, flags)), \
121         [tif_fork] "i" (_TIF_FORK), \
122         [thread_info] "i" (offsetof(struct task_struct, stack)), \
123         [current_task] "m" (current_task) \
124         : \
125         : "memory", "cc" __EXTRA_CLOBBER)
126
127 #endif /* CONFIG_X86_32 */
128
129 #endif /* _ASM_X86_SWITCH_TO_H */
130
```

switching  
kernel stack  
happens in  
here







# Context switch / switch kernel stack

```
← → ↻ lxr.linux.no/linux+v3.6.3/arch/x86/kernel/process_64.c#L269
_260 *
_261 * This could still be optimized:
_262 * - fold all the options into a flag word and test it with a single test.
_263 * - could test fs/gs bitsliced
_264 *
_265 * Kprobes not supported here. Set the probe on schedule instead.
_266 * Function graph tracer not supported too.
_267 */
_268 notrace_funcgraph struct task_struct *
_269 __switch_to(struct task_struct *prev_p, struct task_struct *next_p)
_270 {
_271     struct thread_struct *prev = &prev_p->thread;
_272     struct thread_struct *next = &next_p->thread;
_273     int cpu = smp_processor_id();
_274     struct tss_struct *tss = &per_cpu(init_tss, cpu);
_275     unsigned fsindex, gsindex;
_276     fpu_switch_t fpu;
_277
_278     fpu = switch_fpu_prepare(prev_p, next_p, cpu);
_279
_280     /*
_281     * Reload esp0, LDT and the page table pointer:
_282     */
_283     load_sp0(tss, next); ←
_284
_285     /*
_286     * Switch DS and ES.
_287     * This won't pick up thread selector changes, but I guess that is ok.
_288     */
_289     savesegment(es, prev->es);
_290     if (unlikely(next->es != prev->es))
_291         loadsegment(es, next->es);
_292
_293     savesegment(ds, prev->ds);
_294     if (unlikely(next->ds != prev->ds))
_295         loadsegment(ds, next->ds);
_296
_297     /* We must save %fs and %gs before load_TLS() because
_298     * %fs and %gs may be cleared by load_TLS().
_299     *
_300     * (e.g. xen_load_tls())
_301     */
_302
_303     savesegment(fs, fsindex);
_304     savesegment(gs, gsindex);
_305
_306     load_TLS(next, cpu);
_307
```

Switch kernel  
stack





# Scheduling Criteria

---

- ☐ **Throughput** – # of processes that complete their execution per time unit
  - ☐ Higher is better
- ☐ **Turnaround time** – amount of time to execute a particular process
  - ☐ Lower is better
- ☐ **Response time** – time from request to first response
  - ☐ E.g. mouse clicking on the menu bar to the showing of the menu
  - ☐ Lower is better
- ☐ Above criteria are affected by secondary criteria
  - ☐ **CPU utilization** – keep the CPU as busy as possible
  - ☐ **Waiting time** – amount of time a process has been waiting in the ready queue






# First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- ☐ Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



- ☐ Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- ☐ Average waiting time:  $(0 + 24 + 27)/3 = 17$
- ☐ Throughput: 3 processes / 30 seconds = 0.1 processes / sec
- ☐ Turnaround Time:  $P_1$ : 24,  $P_2$ : 27,  $P_3$ : 30
  - ☐ Avg. TT:  $(24+27+30)/3 = 27$





# FCFS Scheduling (Cont)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

☐ The Gantt chart for the schedule is:



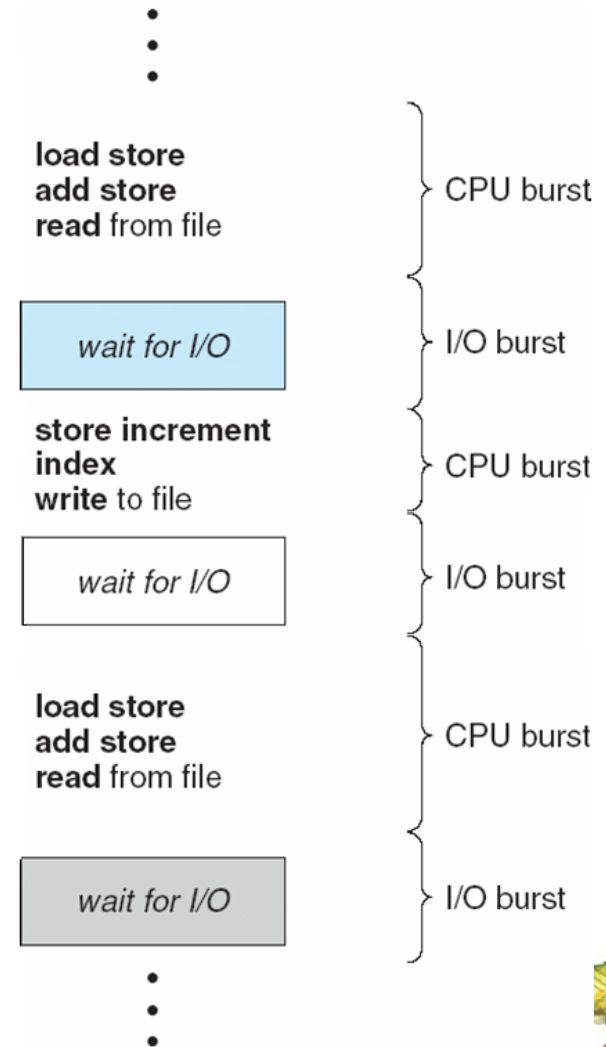

- ☐ Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- ☐ Average waiting time:  $(6 + 0 + 3)/3 = 3$
- ☐ Throughput:  $3 / 30 = 0.1$  processes / sec
- ☐ Turnaround time: Time:  $P_1: 30$ ,  $P_2: 3$ ,  $P_3: 6$ 
  - ☐ Avg. TT:  $(30+3+6)/3 = 13$
- ☐ Much better than previous case
- ☐ Scheduling algorithm can reduce TT
  - ☐ Minimize waiting time to minimize TT
- ☐ What about throughput?





# Alternating Sequence of CPU And I/O Bursts


- CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution, I/O wait, and event wait
- An I/O device can be considered as a special purpose CPU
- Goal: keep all CPUs and all I/O devices busy





# Histogram of CPU-burst Times

Many short CPU bursts and a few long bursts



What does this mean for FCFS?





# FCFS Convoy effect

---

- ☐ CPU bound jobs will hold CPU until exit or I/O
  - ☐ I/O rare for CPU-bound thread
  - ☐ Long periods where no I/O requests issued, and CPU held
  - ☐ => poor I/O device utilization
- ☐ Example: one CPU-bound job, many I/O bound
  - ☐ CPU bound runs (I/O device idle)
  - ☐ CPU bound blocks
  - ☐ I/O bound job(s) run, quickly block on I/O
  - ☐ CPU bound runs again
  - ☐ I/O completes
  - ☐ CPU bound still runs while I/O device idle (continue?)





# Shortest-Job-First (SJF) Scheduling

---

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Process with shortest burst goes next
  - If tie then use FCFS to break tie
- SJF is optimal – gives minimum average waiting time for a given set of processes
  - The difficulty is knowing the length of the next CPU request
- Two schemes:
  - *Non-preemptive* – once CPU assigned, process not preempted until its CPU burst completes
  - *Preemptive* – if a process with CPU burst less than remaining time of current, preempt
    - ▶ *Shortest-Remaining-Time-First (SRTF)*







# Example of Non-Preemptive SJF


- $T = 0$ :  $RQ = \{P_1\}$   
*Select  $P_1$*
- $T = 2$ :  $RQ = \{P_2\}$   
*No-Preemption*
- $T = 4$ :  $RQ = \{P_3, P_2\}$   
*No-Preemption*
- $T = 5$ :  $RQ = \{P_3, P_2, P_4\}$   
*No-Preemption*
- $T = 7$ :  $RQ = \{P_3, P_2, P_4\}$   
 *$P_1$  completes, Select  $P_3$*
- $T = 8$ :  $RQ = \{P_2, P_4\}$   
 *$P_3$  completes, Select  $P_2$*
- $T = 12$ :  $RQ = \{P_4\}$   
 *$P_2$  completes, Select  $P_4$*
- $T = 16$ :  $RQ = \{\}$   
 *$P_4$  completes*

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- Average Waiting Time:  

$$[0 + (8 - 2) + (7 - 4) + (12 - 5)]/4 =$$

$$[6 + 3 + 7]/4 = 4$$






# Example of Preemptive SJF

- $T = 0$ :  $RQ = \{P_1\}$   
*Select  $P_1$*
- $T = 2$ :  $RQ = \{P_2\}$   
*preempt  $P_1$ , Select  $P_2$*
- $T = 4$ :  $RQ = \{P_3, P_1\}$   
*preempt  $P_2$ , Select  $P_3$*
- $T = 5$ :  $RQ = \{P_2, P_4, P_1\}$   
 *$P_3$  completes, Select  $P_2$*
- $T = 7$ :  $RQ = \{P_4, P_1\}$   
 *$P_2$  completes, Select  $P_4$*
- $T = 11$ :  $RQ = \{P_1\}$   
 *$P_4$  completes, Select  $P_1$*
- $T = 16$ :  $RQ = \{\}$   
 *$P_1$  completes*

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- Average Waiting Time:  

$$\frac{[(11-2) + (5-4) + (0) + (7-5)]}{4} = \frac{[9 + 1 + 0 + 2]}{4} = 3$$





# Determining Length of Next CPU Burst

---

- ☐ Can only estimate the length
- ☐ Can be done by using the length of previous CPU bursts, using exponential averaging


1.  $t_n$  = actual length of  $n^{th}$  CPU burst
2.  $\tau_{n+1}$  = predicted value for the next CPU burst
3.  $\alpha, 0 \leq \alpha \leq 1$
4. Define :  $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n.$





# SJF Optimality

Gain of “short”




Proof that the SJF algorithm is optimal

**Gain of short > Loss of long**





# Prediction of the Length of the Next CPU Burst



CPU burst ( $t_i$ )	6	4	6	4	13	13	13	...	
"guess" ( $\tau_i$ )	10	8	6	6	5	9	11	12	...





# Examples of Exponential Averaging

- ☐  $\alpha = 0$ 
  - ☐  $\tau_{n+1} = \tau_n$
  - ☐ Recent history does not count
- ☐  $\alpha = 1$ 
  - ☐  $\tau_{n+1} = \alpha t_n$
  - ☐ Only the actual last CPU burst counts

- ☐ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- ☐ Since both  $\alpha$  and  $(1 - \alpha)$  are less than or equal to 1, each successive term has less weight than its predecessor





# Priority Scheduling

---

- ☐ A priority number (integer) is associated with each process
- ☐ The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - ☐ Preemptive
  - ☐ Non-preemptive
- ☐ SJF is a priority scheduling where priority is the predicted next CPU burst time
- ☐ Problem  $\equiv$  **Starvation** – low priority processes may never execute
- ☐ Solution  $\equiv$  **Aging** – as time progresses increase the priority of the process





# Round Robin (RR)

---

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow$  processor sharing (appears as dedicated processor with speed  $1/n$  actual)
  - $q$  must be large with respect to context switch, otherwise overhead is too high








# Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

□ The Gantt chart is:





□ Typically, higher average turnaround than SJF, but better *response*






# Time Quantum and Context Switch Time





# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7





# Multilevel Queue

---

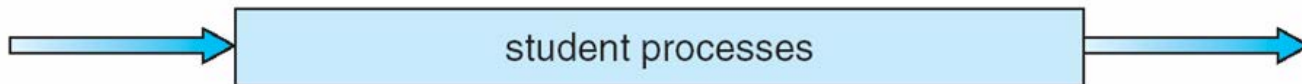

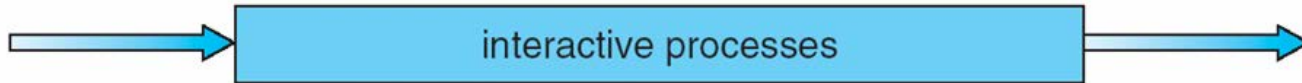
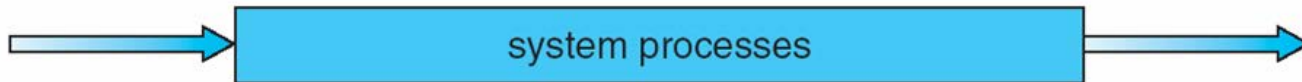
- ☐ Ready queue is partitioned into separate queues:
  - foreground (interactive)
  - background (batch)
- ☐ Each queue has its own scheduling algorithm
  - ☐ foreground – RR
  - ☐ background – FCFS
- ☐ Scheduling must be done between the queues
  - ☐ Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
  - ☐ Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
  - ☐ 20% to background in FCFS





# Multilevel Queue Scheduling

highest priority



lowest priority





# Multilevel Feedback Queue

---

- ☐ A process can move between the various queues; aging can be implemented this way
- ☐ Multilevel-feedback-queue scheduler defined by the following parameters:
  - ☐ number of queues
  - ☐ scheduling algorithms for each queue
  - ☐ method used to determine when to upgrade a process
  - ☐ method used to determine when to demote a process
  - ☐ method used to determine which queue a process will enter when that process needs service






# Multilevel Feedback Queues

1. Each priority level has a ready queue, and a time quantum
2. A new process is positioned at the end of the top-level queue
3. At some stage the process reaches the head of the queue and is assigned the CPU
4. If the process is completed it leaves the system
5. If the process voluntarily relinquishes control it leaves the queuing network, and when the process becomes ready again it enters the system on the same queue level.
6. If the process uses all the quantum time, it is pre-empted and positioned at the end of the next lower level queue.
7. At the base level queue the processes circulate in round robin fashion until they complete and leave the system.
8. If a process blocks for I/O, it is 'promoted' one level, and placed at the end of the next-higher queue. This allows I/O bound processes to be favored by the scheduler and allows processes to 'escape' the base level queue.

Multiple FIFO queues





# Multilevel Feedback Queues

---

- I/O bound processes tend to congregate in higher-level queues
  - This implies greater device utilization
- CPU-bound processes will sink deeper into the queues
- Quantum in top queue should be large enough to satisfy majority of I/O-bound processes
- Can assign a process a lower priority by starting it at a lower-level queue
- Can raise priority by moving process to a higher queue, thus can use in conjunction with aging







# Thread Scheduling

---

- ☐ Distinction between user-level and kernel-level threads
- ☐ Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
  - ☐ Known as **process-contention scope (PCS)** since scheduling competition is within the process
- ☐ Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in system





# Pthread Scheduling

---

- API allows specifying either PCS or SCS during thread creation
  - PTHREAD SCOPE PROCESS schedules threads using PCS scheduling
  - PTHREAD SCOPE SYSTEM schedules threads using SCS scheduling.





# Pthread Scheduling API

---

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[])
{
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_t init(&attr);
    /* set the scheduling algorithm to PROCESS or SYSTEM */
    pthread_attr_t setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    /* set the scheduling policy - FIFO, RT, or OTHER */
    pthread_attr_t setschedpolicy(&attr, SCHED_OTHER);
    /* create the threads */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
}
```





# Pthread Scheduling API

---

```
/* now join on each thread */
for (i = 0; i < NUM THREADS; i++)
    pthread join(tid[i], NULL);
}

/* Each thread will begin control in this function */
void *runner(void *param)
{
    printf("I am a thread\n");
    pthread exit(0);
}
```





# Multiple-Processor Scheduling


---

- ☐ CPU scheduling more complex when multiple CPUs are available
- ☐ **Homogeneous processors** within a multiprocessor
- ☐ **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
- ☐ **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
- ☐ **Processor affinity** – process has affinity for processor on which it is currently running
  - ☐ **soft affinity**
  - ☐ **hard affinity**





# NUMA and CPU Scheduling





# Multicore Processors

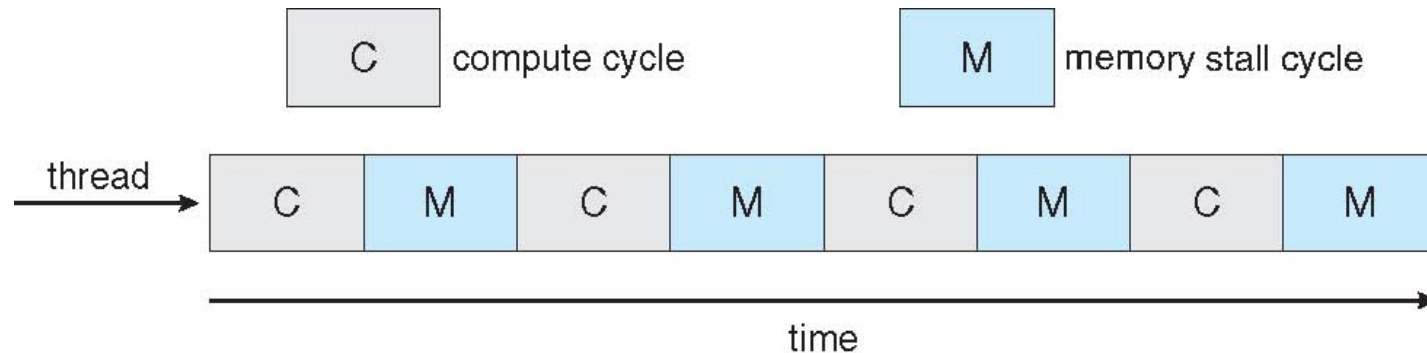
---

- ☐ Recent trend to place multiple processor cores on same physical chip
- ☐ Faster and consume less power
- ☐ Multiple threads per core also growing
  - ☐ Takes advantage of memory stall to make progress on another thread while memory retrieve happens





# Multithreaded Multicore System







# Operating System Examples

---

- ☐ Solaris scheduling
- ☐ Windows XP scheduling
- ☐ Linux scheduling






# Solaris Dispatch Table

priority	time quantum	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56
50	40	40	58
55	40	45	58
59	20	49	59





# Solaris Scheduling





# Windows XP Priorities

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1





# Linux Scheduling

---

- ☐ Constant order  $O(1)$  scheduling time
- ☐ Two priority ranges: time-sharing and real-time
- ☐ **Real-time** range from 0 to 99 and **nice** value from 100 to 140
- ☐ (figure 5.15)





# Priorities and Time-slice length


<u>numeric priority</u>	<u>relative priority</u>		<u>time quantum</u>
0	highest	real-time tasks	200 ms
•			
•			
•			
99			
100		other tasks	
•			
•			
•			
140	lowest		10 ms






# List of Tasks Indexed According to Priorities

## active array



## expired array





# Algorithm Evaluation

---

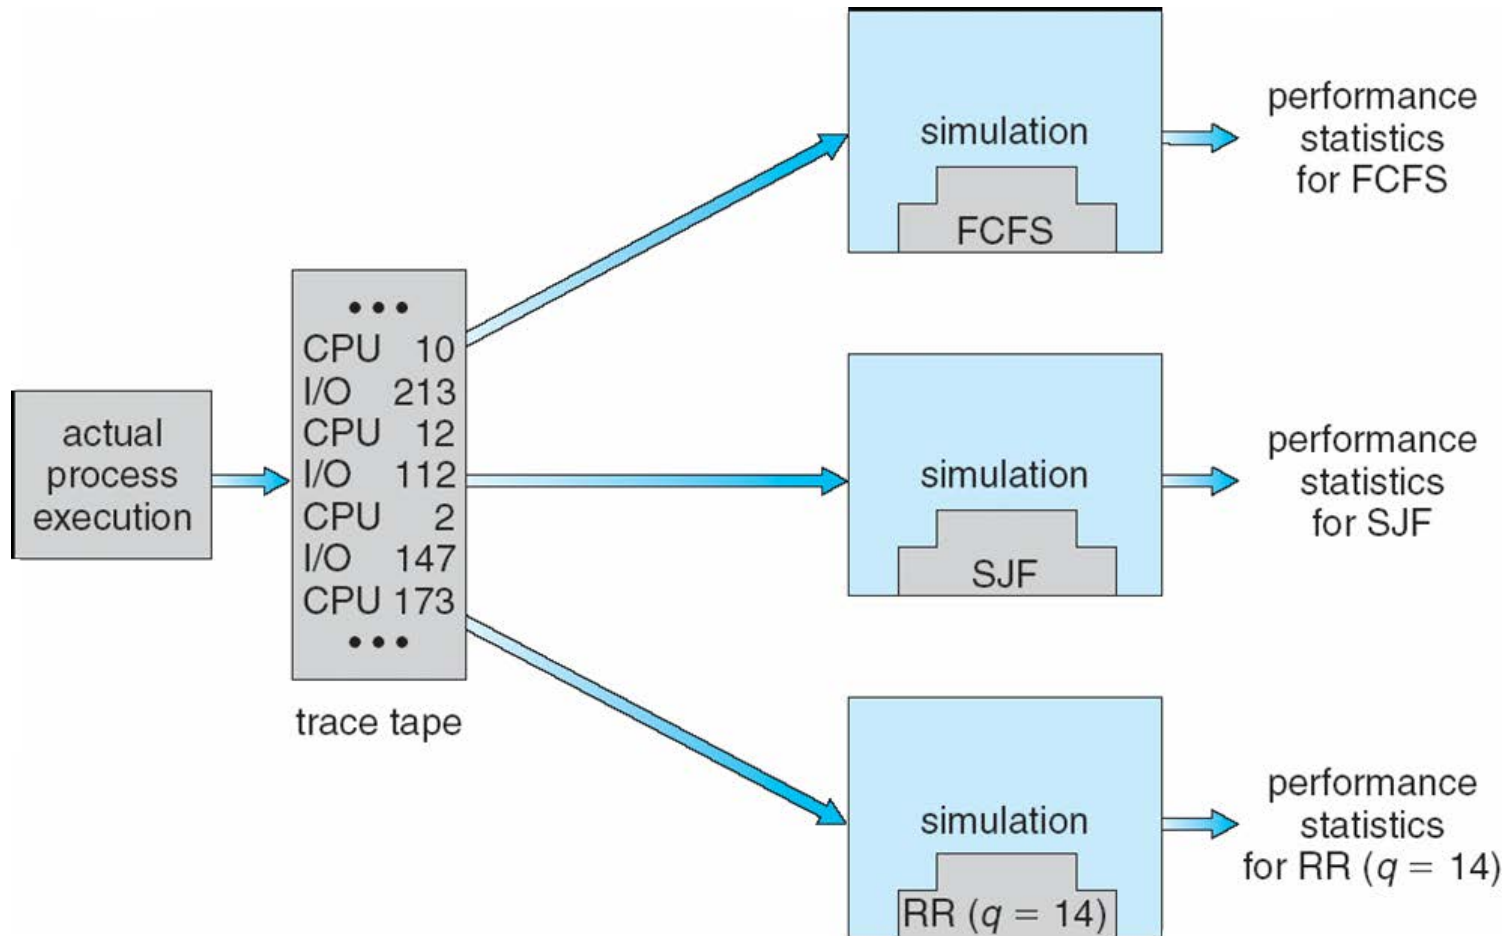
- ☐ Deterministic modeling – takes a particular predetermined workload and defines the performance of each algorithm for that workload
- ☐ Queueing models
- ☐ Implementation








# Evaluation of CPU schedulers by Simulation




# End of Chapter 5

---





## 5.08





# In-5.7


---





## In-5.8


---





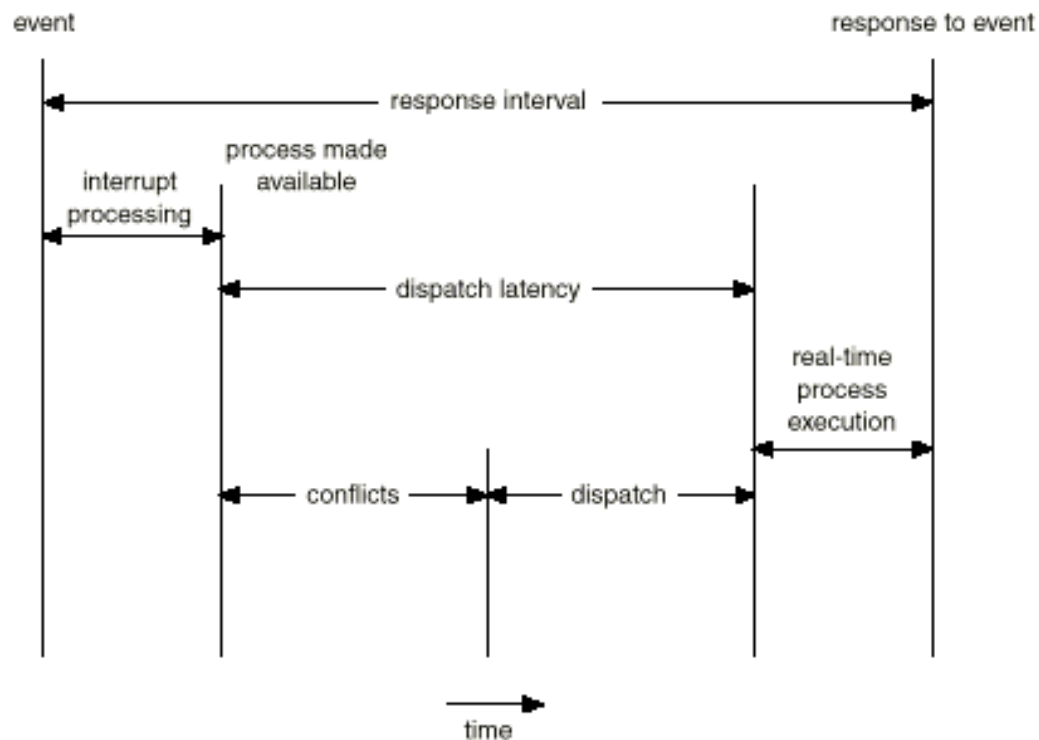
# In-5.9

---





# Dispatch Latency





# Java Thread Scheduling

---

- ☐ JVM Uses a Preemptive, Priority-Based Scheduling Algorithm
- ☐ FIFO Queue is Used if There Are Multiple Threads With the Same Priority







# Java Thread Scheduling (cont)

---

JVM Schedules a Thread to Run When:

1. The Currently Running Thread Exits the Runnable State
2. A Higher Priority Thread Enters the Runnable State

\* Note – the JVM Does Not Specify Whether Threads are Time-Sliced or Not





# Time-Slicing

---

Since the JVM Doesn't Ensure Time-Slicing, the yield() Method May Be Used:

```
while (true) {  
    // perform CPU-intensive task  
    . . .  
    Thread.yield();  
}
```

This Yields Control to Another Thread of Equal Priority





# Thread Priorities

---

<u>Priority</u>	<u>Comment</u>
Thread.MIN_PRIORITY	Minimum Thread Priority
Thread.MAX_PRIORITY	Maximum Thread Priority
Thread.NORM_PRIORITY	Default Thread Priority

Priorities May Be Set Using setPriority() method:

```
setPriority(Thread.NORM_PRIORITY + 2);
```





# Solaris 2 Scheduling

