


Chapter 13: I/O Systems





Chapter 13: I/O Systems

I/O Hardware

Application I/O Interface

Kernel I/O Subsystem

Transforming I/O Requests to Hardware Operations

Streams

Performance





Objectives

Explore the structure of an operating system's I/O subsystem

Discuss the principles of I/O hardware and its complexity

Provide details of the performance aspects of I/O hardware and software





I/O Hardware

Incredible variety of I/O devices

Common concepts

Port

Bus (daisy chain or shared direct access)

Controller (host adapter)

I/O instructions control devices

Devices have addresses, used by

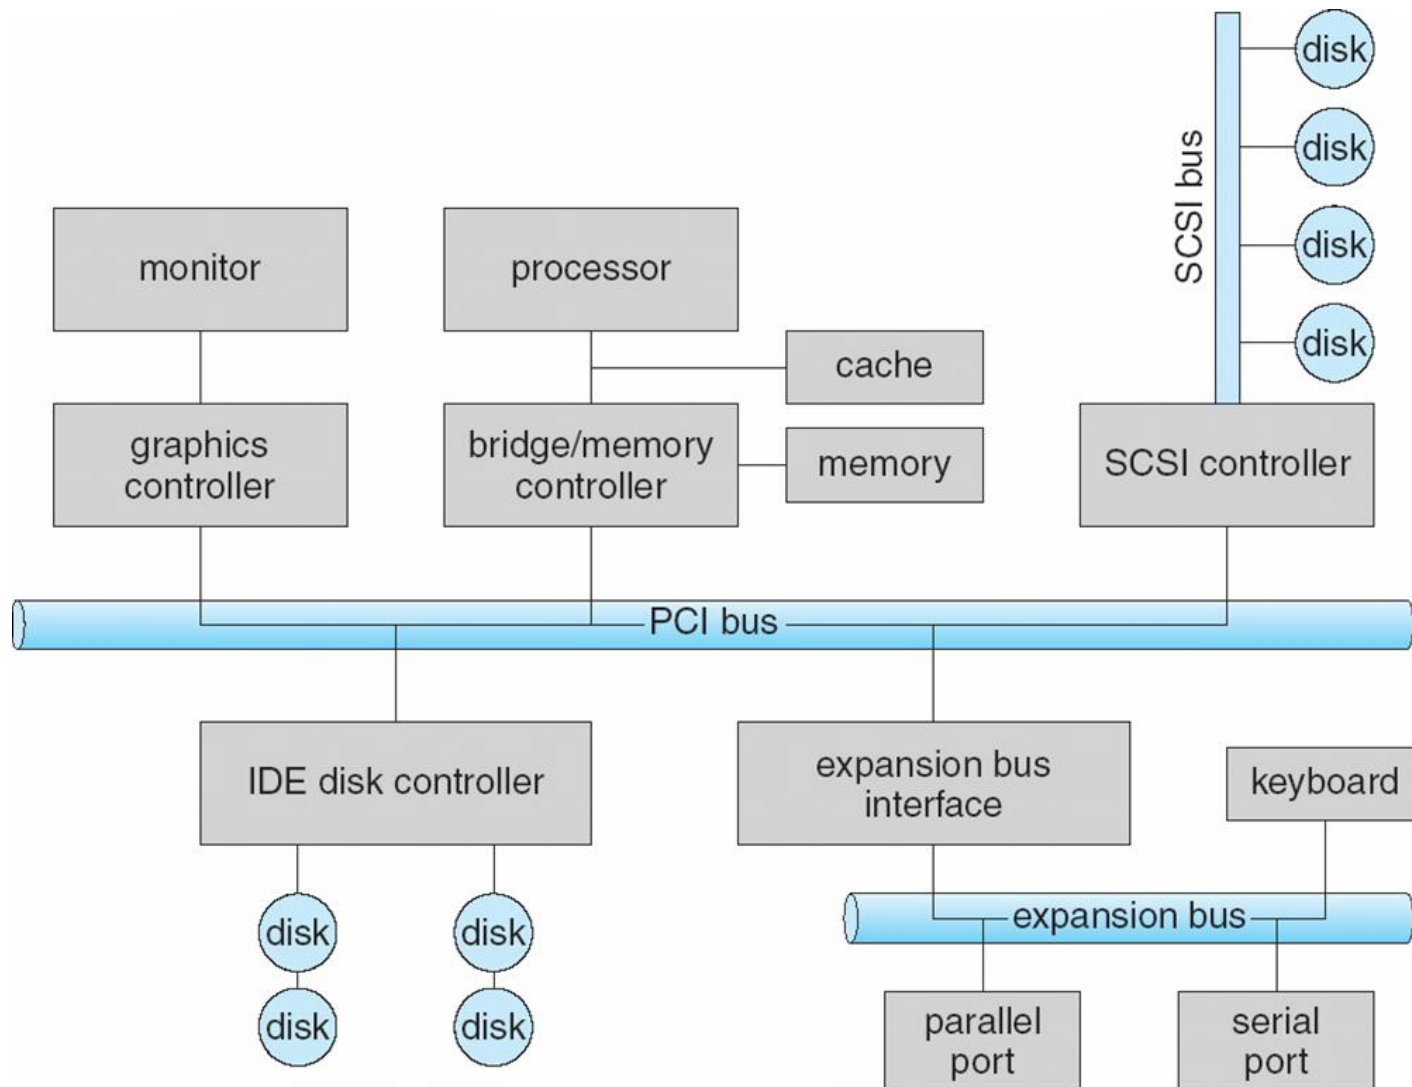
Direct I/O instructions

Memory-mapped I/O





A Typical PC Bus Structure





Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)





Polling

Determines state of device

command-ready

busy

Error

Busy-wait cycle to wait for I/O from device





Interrupts

CPU Interrupt-request line triggered by I/O device

Interrupt handler receives interrupts

Maskable to ignore or delay some interrupts

Interrupt vector to dispatch interrupt to correct handler

Based on priority


Some **nonmaskable**

Interrupt mechanism also used for exceptions





Interrupt-Driven I/O Cycle





Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts






Direct Memory Access


Used to avoid **programmed I/O** for large data movement

```
/*
 * example.c: very simple example of port I/O
 *
 * This code does nothing useful, just a port write, a pause,
 * and a port read. Compile with `gcc -O2 -o example example.c',
 * and run as root with `./example'.
 */
#include <stdio.h>
#include <unistd.h>
#include <sys/io.h>
#define BASEPORT 0x378 /* lpl */
int main()
{
    /* Get access to the ports */
    if (ioperm(BASEPORT, 3, 1)) {perror("ioperm"); return 1;}
    /* Set the data signals (D0-7) of the port to all low (0) */
    outb(0, BASEPORT);
    /* Sleep for a while (100 ms) */
    usleep(100000);
    /* Read from the status port (BASE+1) and display the result */
    printf("status: %d\n", inb(BASEPORT + 1));
    /* We don't need the ports anymore */
    if (ioperm(BASEPORT, 3, 0)) {perror("ioperm"); return 1;}
    return 0;
}
/* end of example.c */
```

Requires **DMA** controller

Bypasses CPU to transfer data directly between I/O device and memory



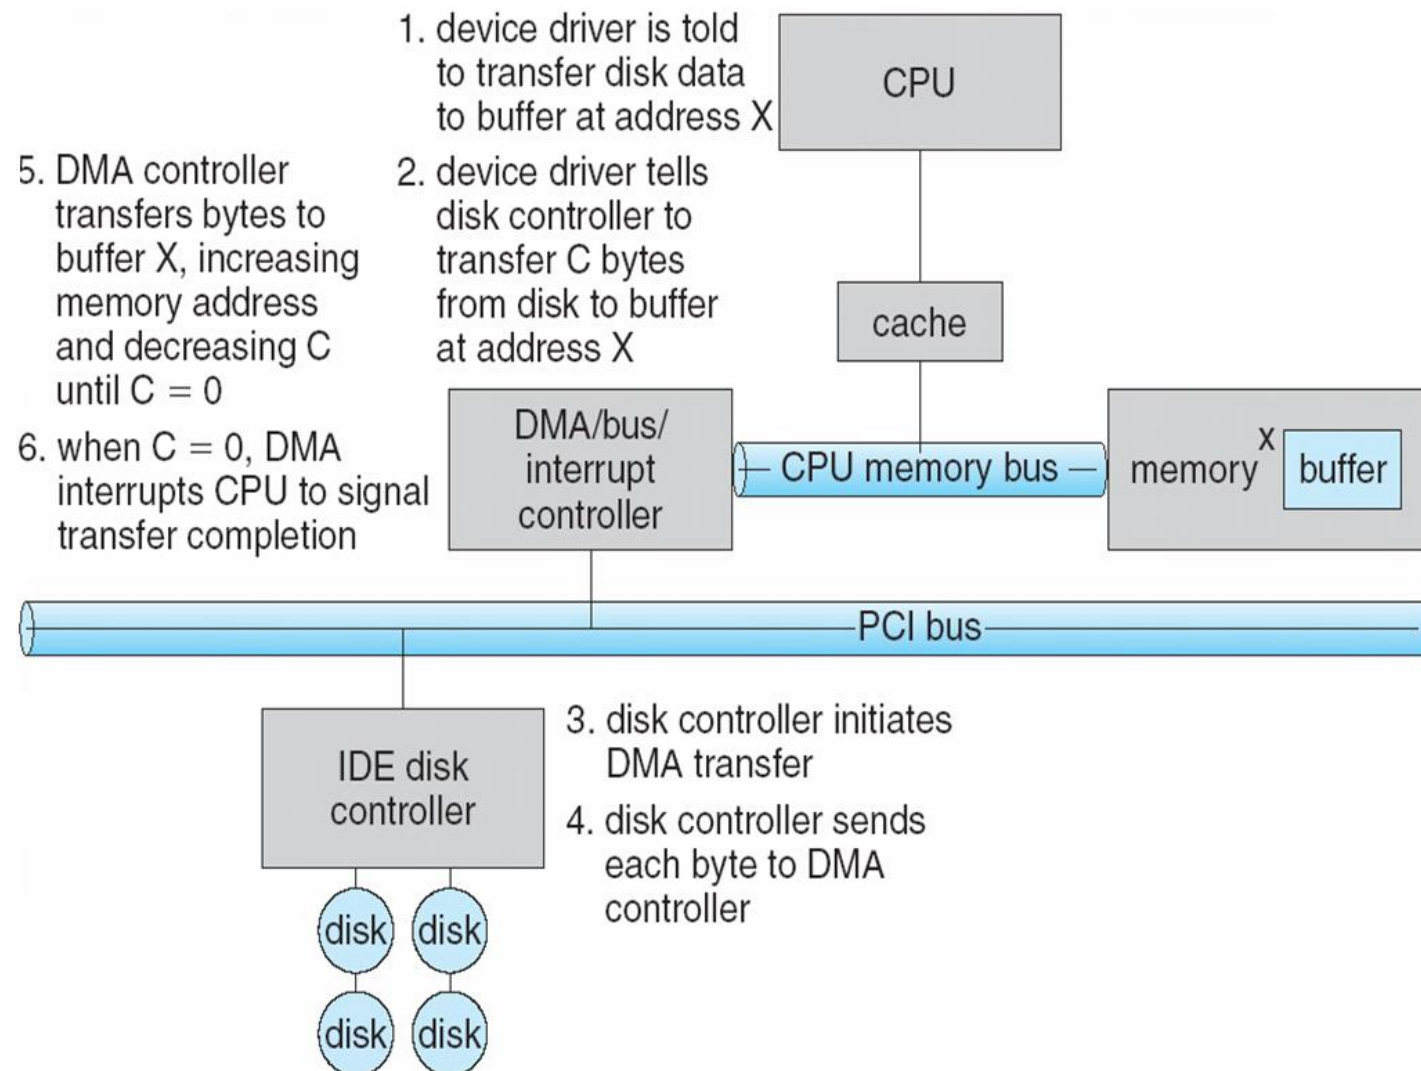


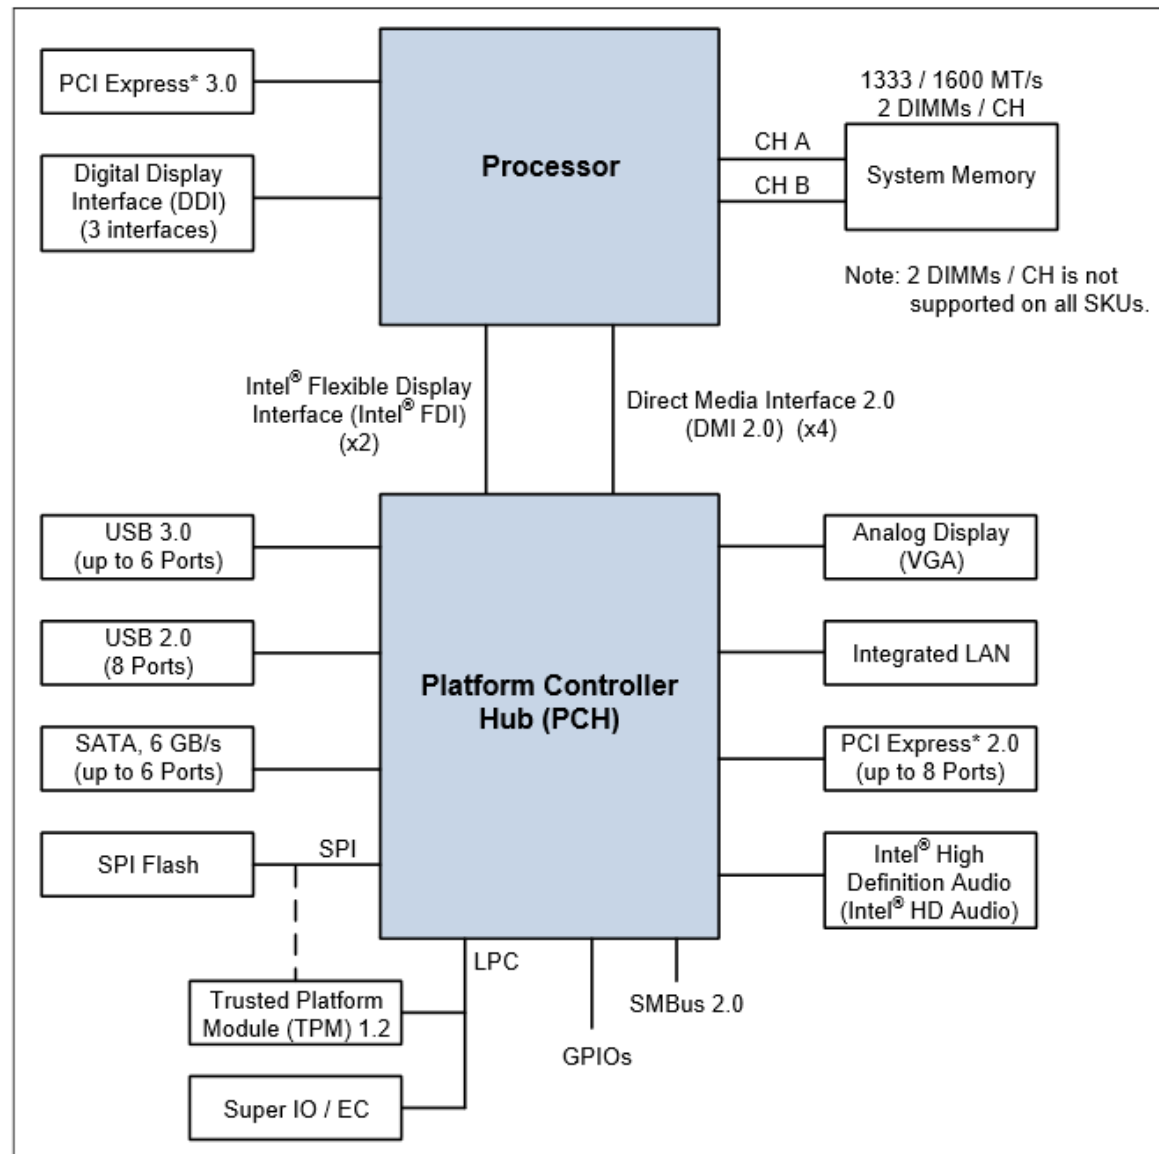
```
hank@hank-VirtualBox: ~  
hank@hank-VirtualBox:~$ sudo cat /proc/ioproports  
[sudo] password for hank:  
0000-001f : dma1  
0020-0021 : pic1  
0040-0043 : timer0  
0050-0053 : timer1  
0060-0060 : keyboard  
0064-0064 : keyboard  
0070-0071 : rtc_cmos  
    0070-0071 : rtc0  
0080-008f : dma page reg  
00a0-00a1 : pic2  
00c0-00df : dma2  
00f0-00ff : fpu  
0170-0177 : 0000:00:01.1  
    0170-0177 : ata_piix  
01f0-01f7 : 0000:00:01.1  
    01f0-01f7 : ata_piix  
0376-0376 : 0000:00:01.1  
    0376-0376 : ata_piix  
03c0-03df : vesafb  
03f6-03f6 : 0000:00:01.1  
    03f6-03f6 : ata_piix  
0cf8-0cff : PCI conf1  
4000-4003 : ACPI PM1a_EVT_BLK  
4004-4005 : ACPI PM1a_CNT_BLK  
4008-400b : ACPI PM_TMR  
4020-4021 : ACPI GPE0_BLK  
d000-d00f : 0000:00:01.1  
    d000-d00f : ata_piix  
d010-d017 : 0000:00:03.0  
    d010-d017 : e1000  
d020-d03f : 0000:00:04.0  
d100-d1ff : 0000:00:05.0  
    d100-d1ff : Intel 82801AA-ICH  
d200-d23f : 0000:00:05.0  
    d200-d23f : Intel 82801AA-ICH  
d240-d247 : 0000:00:0d.0  
    d240-d247 : ahci  
d250-d257 : 0000:00:0d.0  
    d250-d257 : ahci  
d260-d26f : 0000:00:0d.0  
    d260-d26f : ahci  
hank@hank-VirtualBox:~$
```





Six Step Process to Perform DMA Transfer

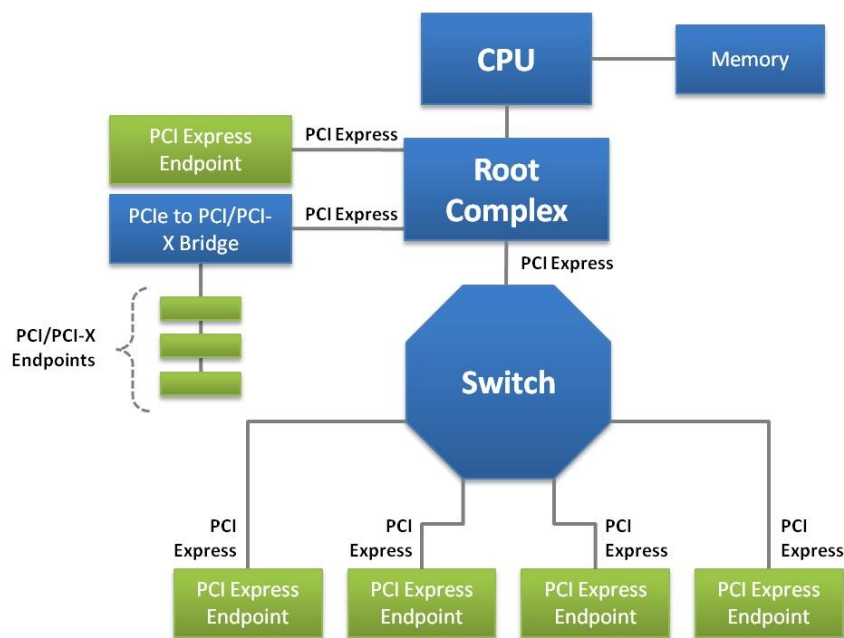






PCI Express

<http://xillybus.com/tutorials/pci-express-tlp-pcie-primer-tutorial-guide-1>





Application I/O Interface

I/O system calls encapsulate device behaviors in generic classes

Device-driver layer hides differences among I/O controllers from kernel

Devices vary in many dimensions

Character-stream or block

Sequential or random-access

Sharable or dedicated

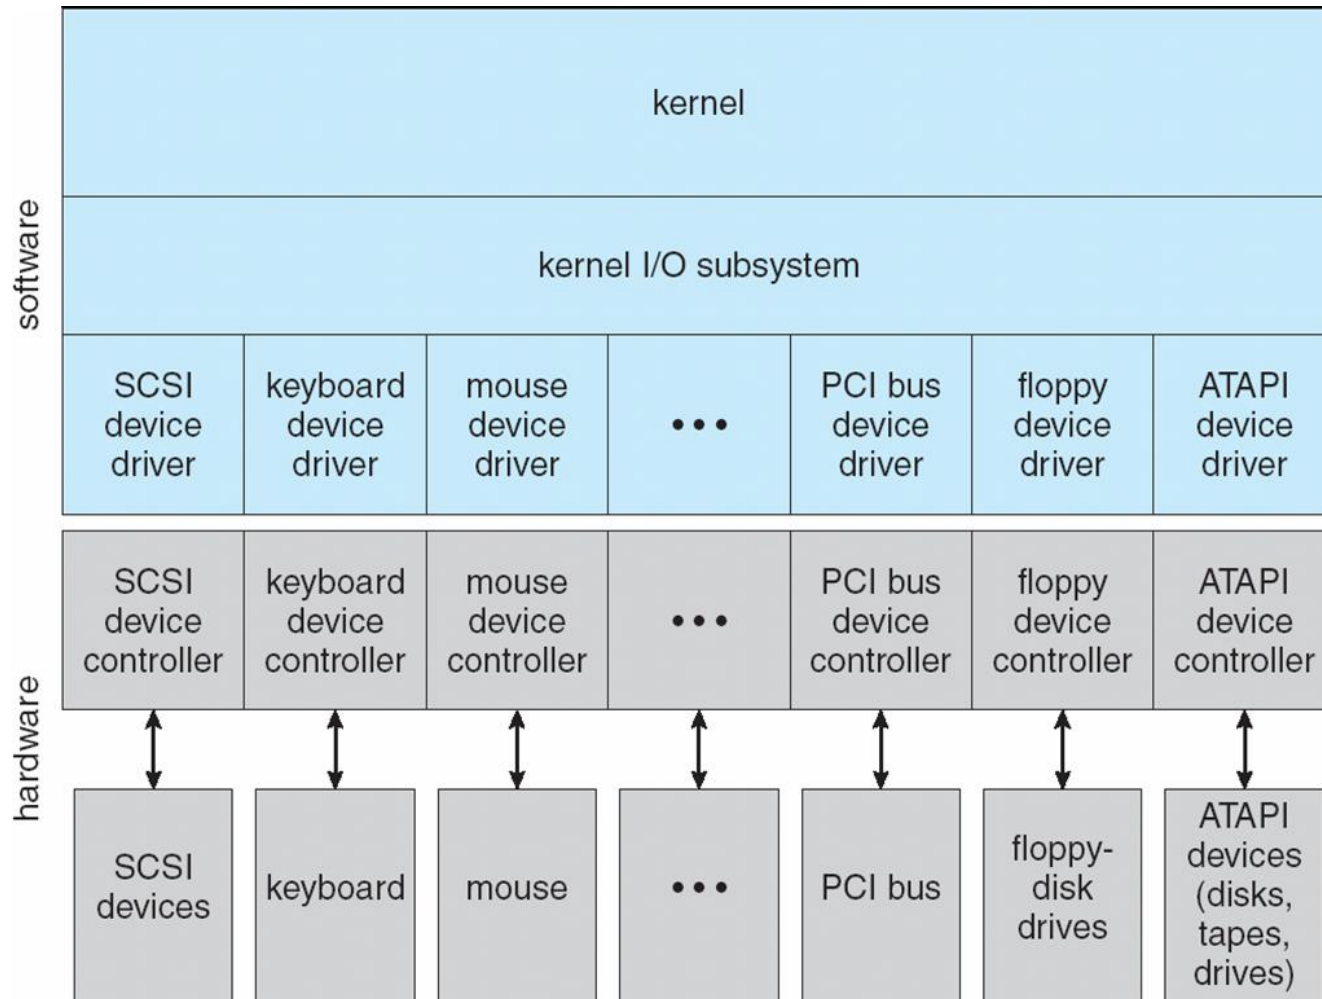
Speed of operation

read-write, read only, or write only





A Kernel I/O Structure





Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk





Block and Character Devices

Block devices include disk drives

Commands include read, write, seek

Raw I/O or file-system access

Memory-mapped file access possible

Character devices include keyboards, mice, serial ports

Commands include `get`, `put`

Libraries layered on top allow line editing





Network Devices

Varying enough from block and character to have own interface

Unix and Windows NT/9x/2000 include socket interface

Separates network protocol from network operation

Includes `select` functionality

Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)





Clocks and Timers

Provide current time, elapsed time, timer

Programmable interval timer used for timings, periodic interrupts

`ioctl` (on UNIX) covers odd aspects of I/O such as clocks and timers





Blocking and Nonblocking I/O

Blocking - process suspended until I/O completed

Easy to use and understand

Insufficient for some needs

Nonblocking - I/O call returns as much as available

User interface, data copy (buffered I/O)

Implemented via multi-threading

Returns quickly with count of bytes read or written

Asynchronous - process runs while I/O executes


Difficult to use

I/O subsystem signals process when I/O completed





Two I/O Methods





Kernel I/O Subsystem

Scheduling

- Some I/O request ordering via per-device queue

- Some OSs try fairness

Buffering - store data in memory while transferring between devices

- To cope with device speed mismatch


- To cope with device transfer size mismatch

- To maintain “copy semantics”







Device-status Table





Sun Enterprise 6000 Device-Transfer Rates





Kernel I/O Subsystem

Caching - fast memory holding copy of data

Always just a copy

Key to performance

Spooling - hold output for a device

If device can serve only one request at a time

i.e., Printing

Device reservation - provides exclusive access to a device

System calls for allocation and deallocation

Watch out for deadlock





Error Handling

OS can recover from disk read, device unavailable, transient write failures

Most return an error number or code when I/O request fails

System error logs hold problem reports





I/O Protection

User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions

All I/O instructions defined to be privileged


I/O must be performed via system calls

- ▶ Memory-mapped and I/O port memory locations must be protected too





Use of a System Call to Perform I/O





Kernel Data Structures

Kernel keeps state info for I/O components, including open file tables, network connections, character device state

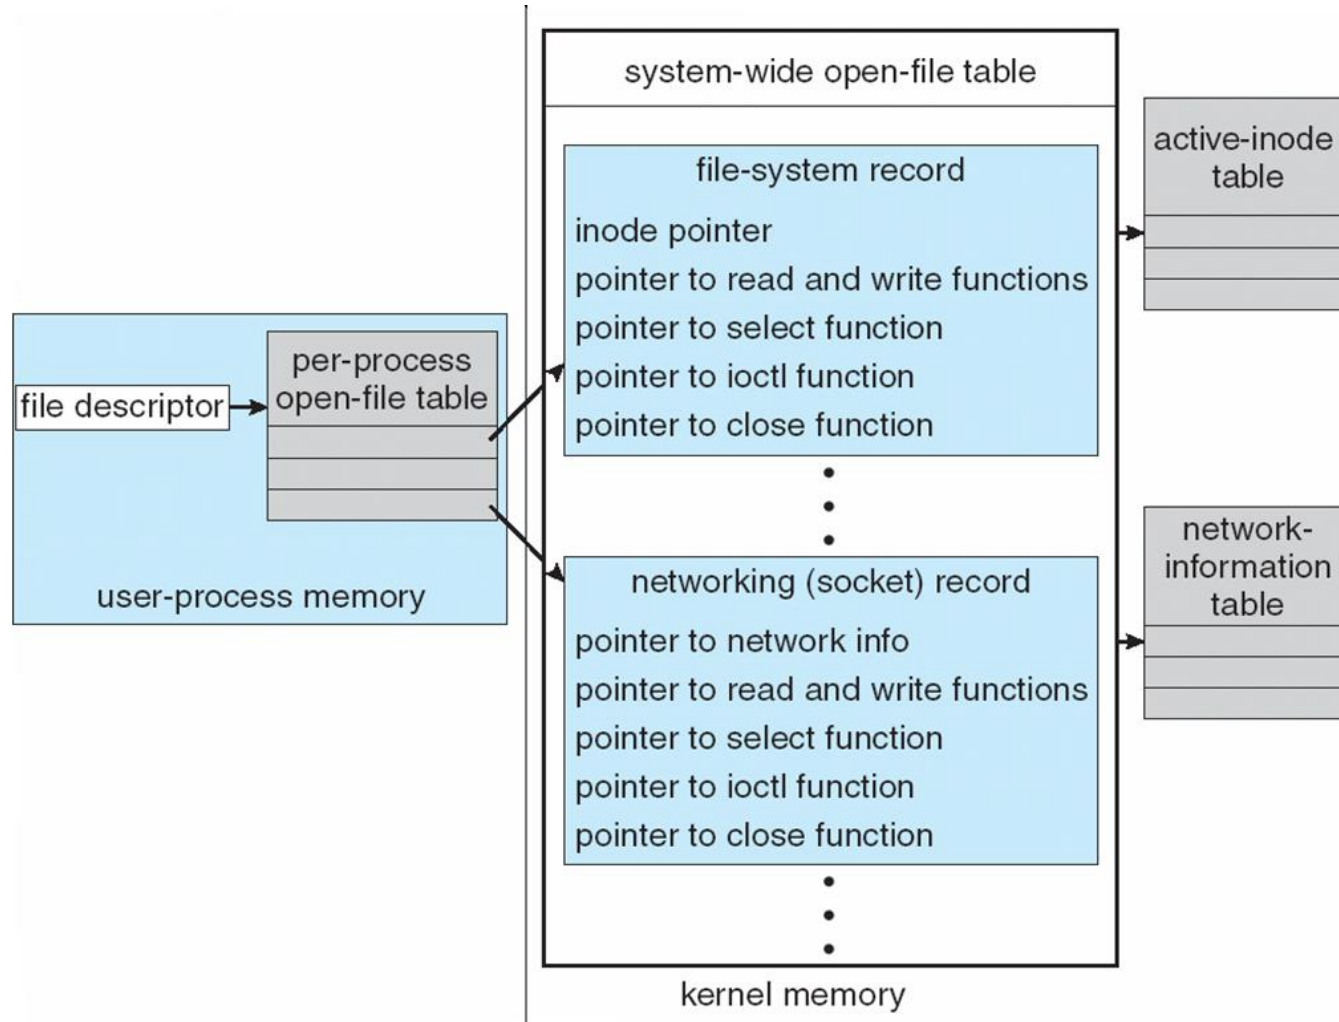
Many, many complex data structures to track buffers, memory allocation, “dirty” blocks

Some use object-oriented methods and message passing to implement I/O





UNIX I/O Kernel Structure





I/O Requests to Hardware Operations

Consider reading a file from disk for a process:

- Determine device holding file

- Translate name to device representation

- Physically read data from disk into buffer


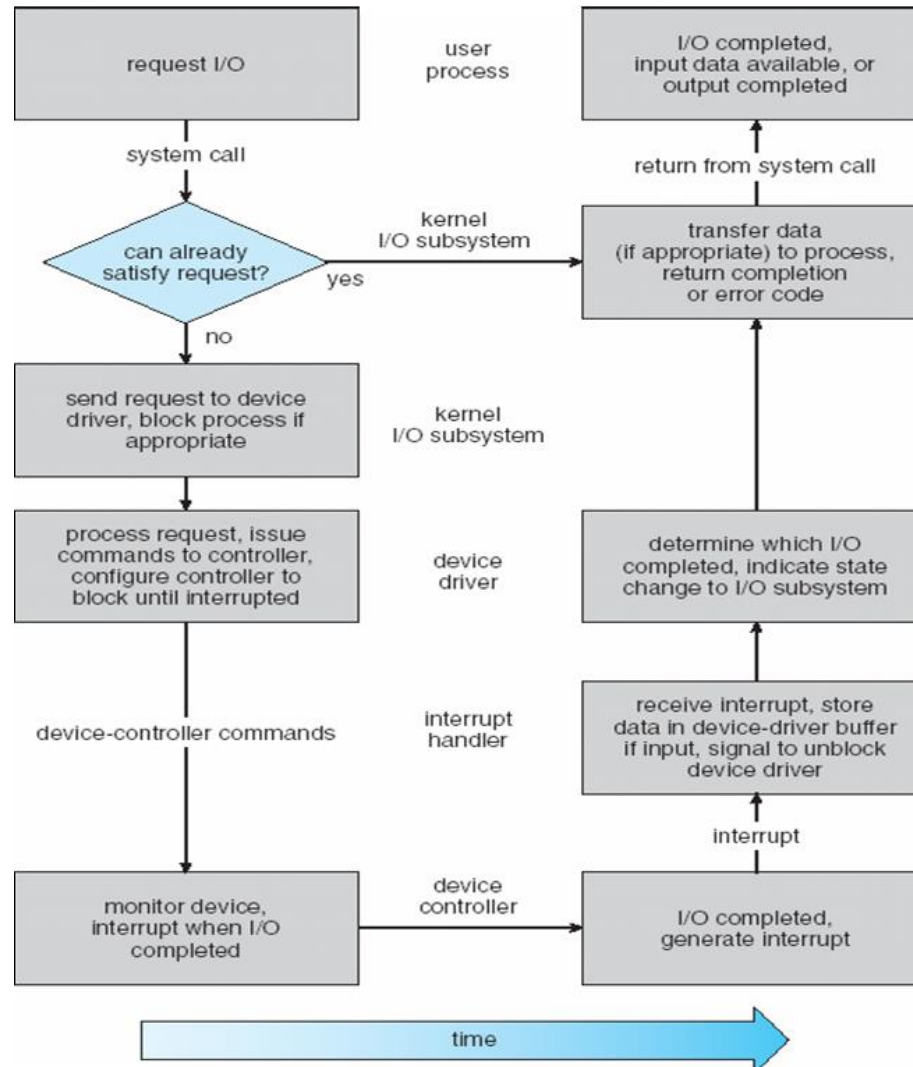
- Make data available to requesting process

- Return control to process





Life Cycle of An I/O Request





STREAMS

STREAM – a full-duplex communication channel between a user-level process and a device in Unix System V and beyond

A STREAM consists of:

- STREAM head interfaces with the user process
- driver end interfaces with the device
- zero or more STREAM modules between them.


Each module contains a **read queue** and a **write queue**

Message passing is used to communicate between queues





The STREAMS Structure





Performance

I/O a major factor in system performance:

- Demands CPU to execute device driver, kernel I/O code

- Context switches due to interrupts


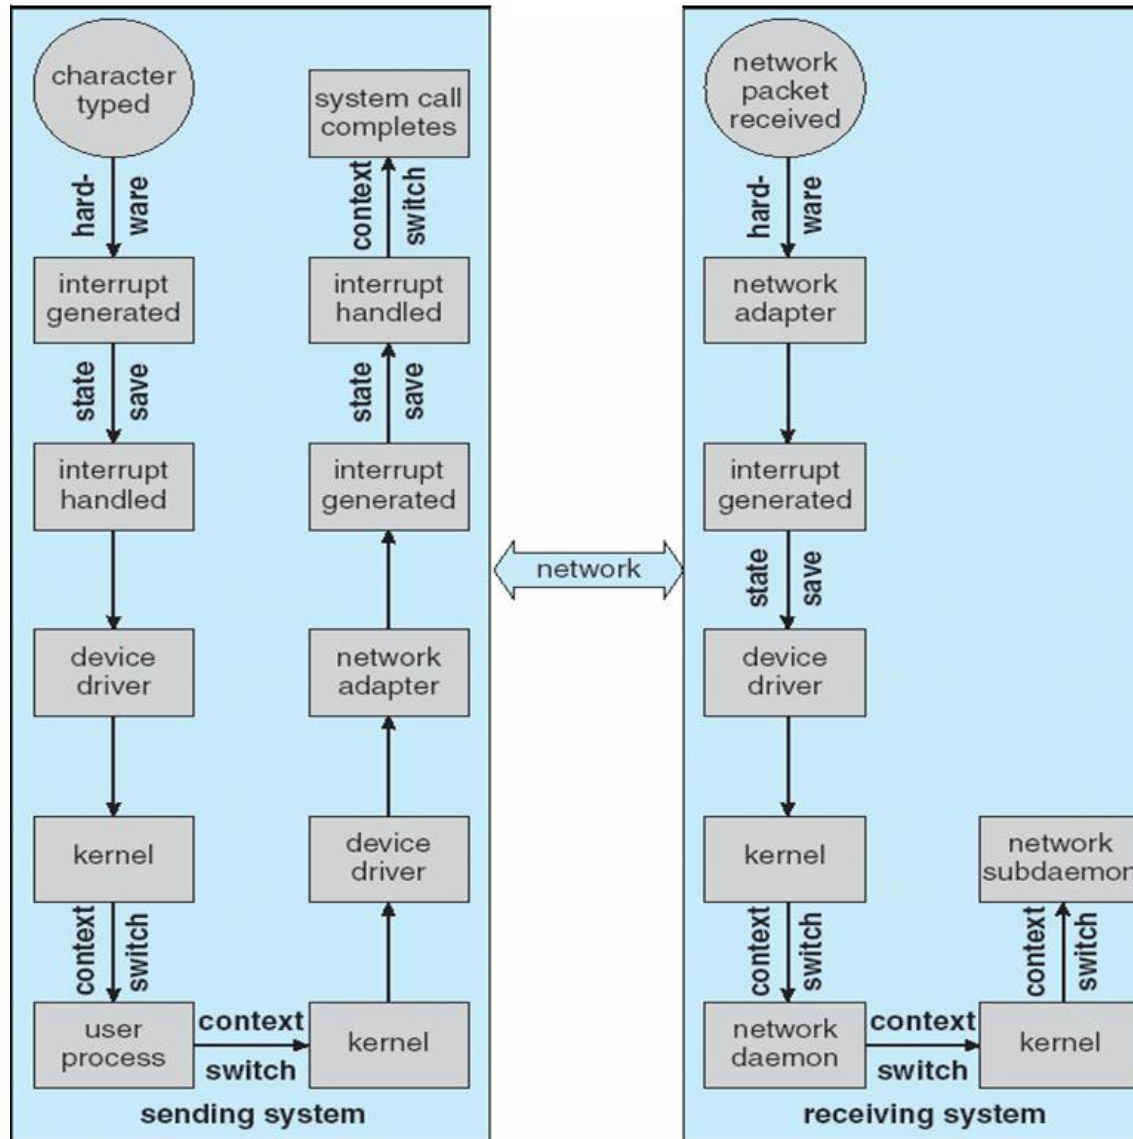
- Data copying

- Network traffic especially stressful





Intercomputer Communications





Improving Performance

Reduce number of context switches

Reduce data copying

Reduce interrupts by using large transfers, smart controllers, polling

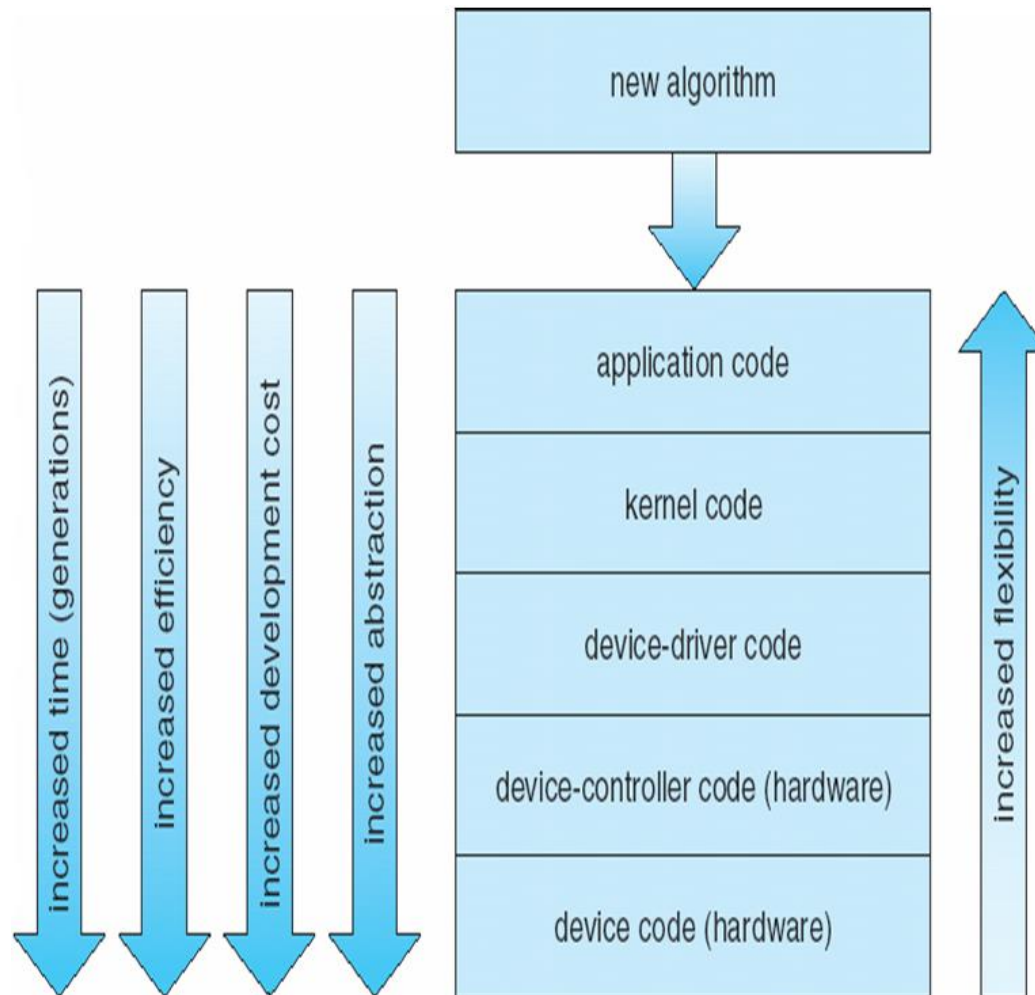
Use DMA

Balance CPU, memory, bus, and I/O performance for highest throughput





Device-Functionality Progression



End of Chapter 13

